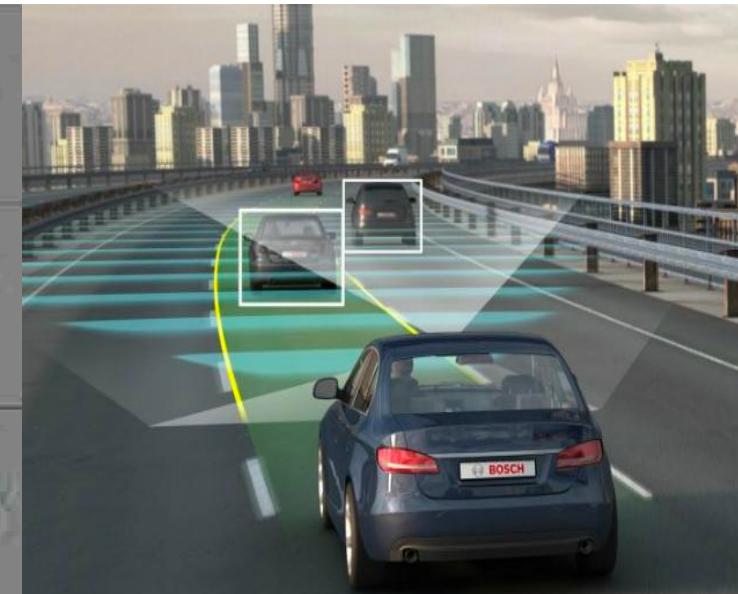
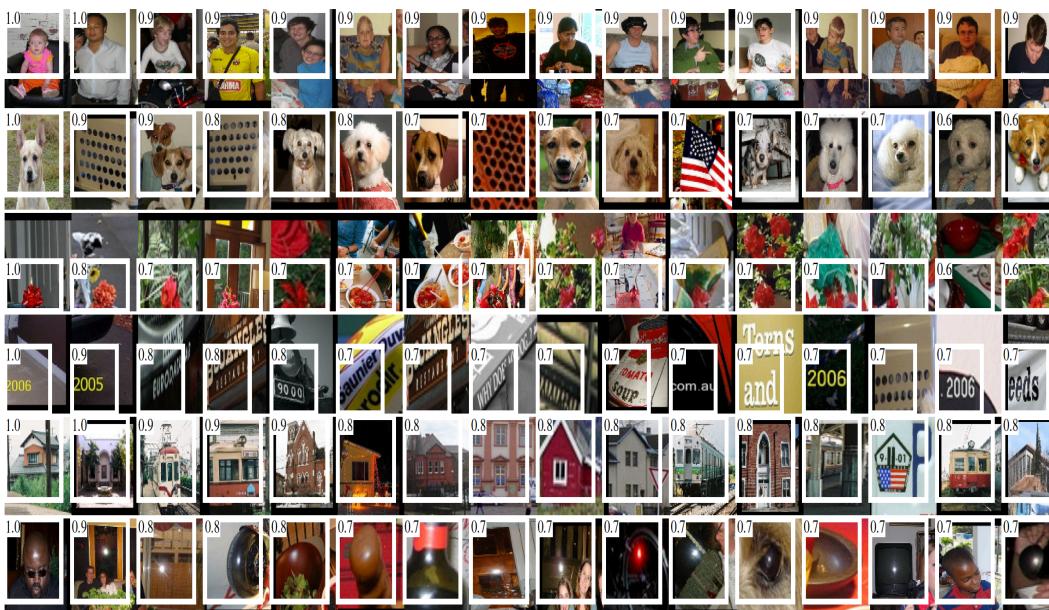


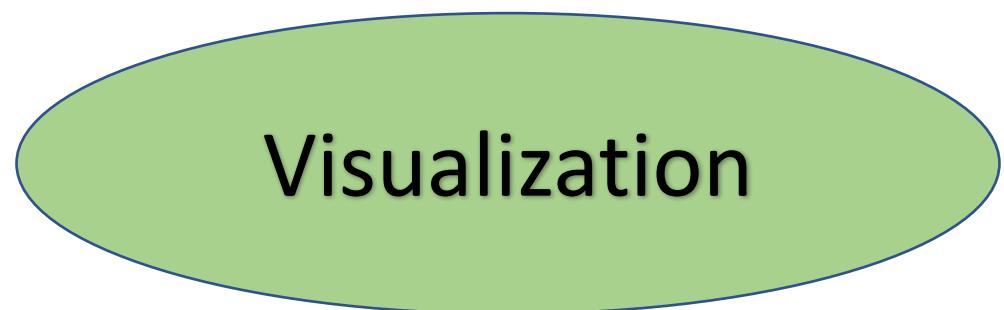
Visualization and Adversarial Examples

Jyoti Aneja, Ralf Gunter Correa Carvalho, Jiahui Yu

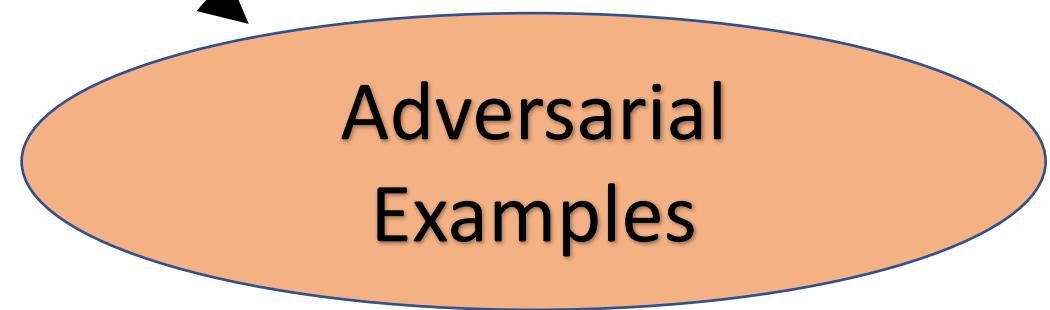
CS-598LAZ



Today's Talk



1. What is Visualization?
2. Visualize patches that maximally activate neurons
3. Visualize the weights
4. Gradient based approaches
5. Optimization based approach



1. Adversarial and Rubbish examples
2. Evolutionary approach
3. Gradient based approaches
4. Adversarial training
5. Transferability
6. Universal Adversarial Perturbations
7. Why are neural networks easily fooled?
8. Proposed Solutions for adversarial attack

Outline - Visualization

- What is Visualization?
- Visualize patches that maximally activate neurons
- Visualize the weights
- Gradient based approaches
- Optimization based approach

What is visualization?

Mapping between a neuron in a layer to the features in the image.

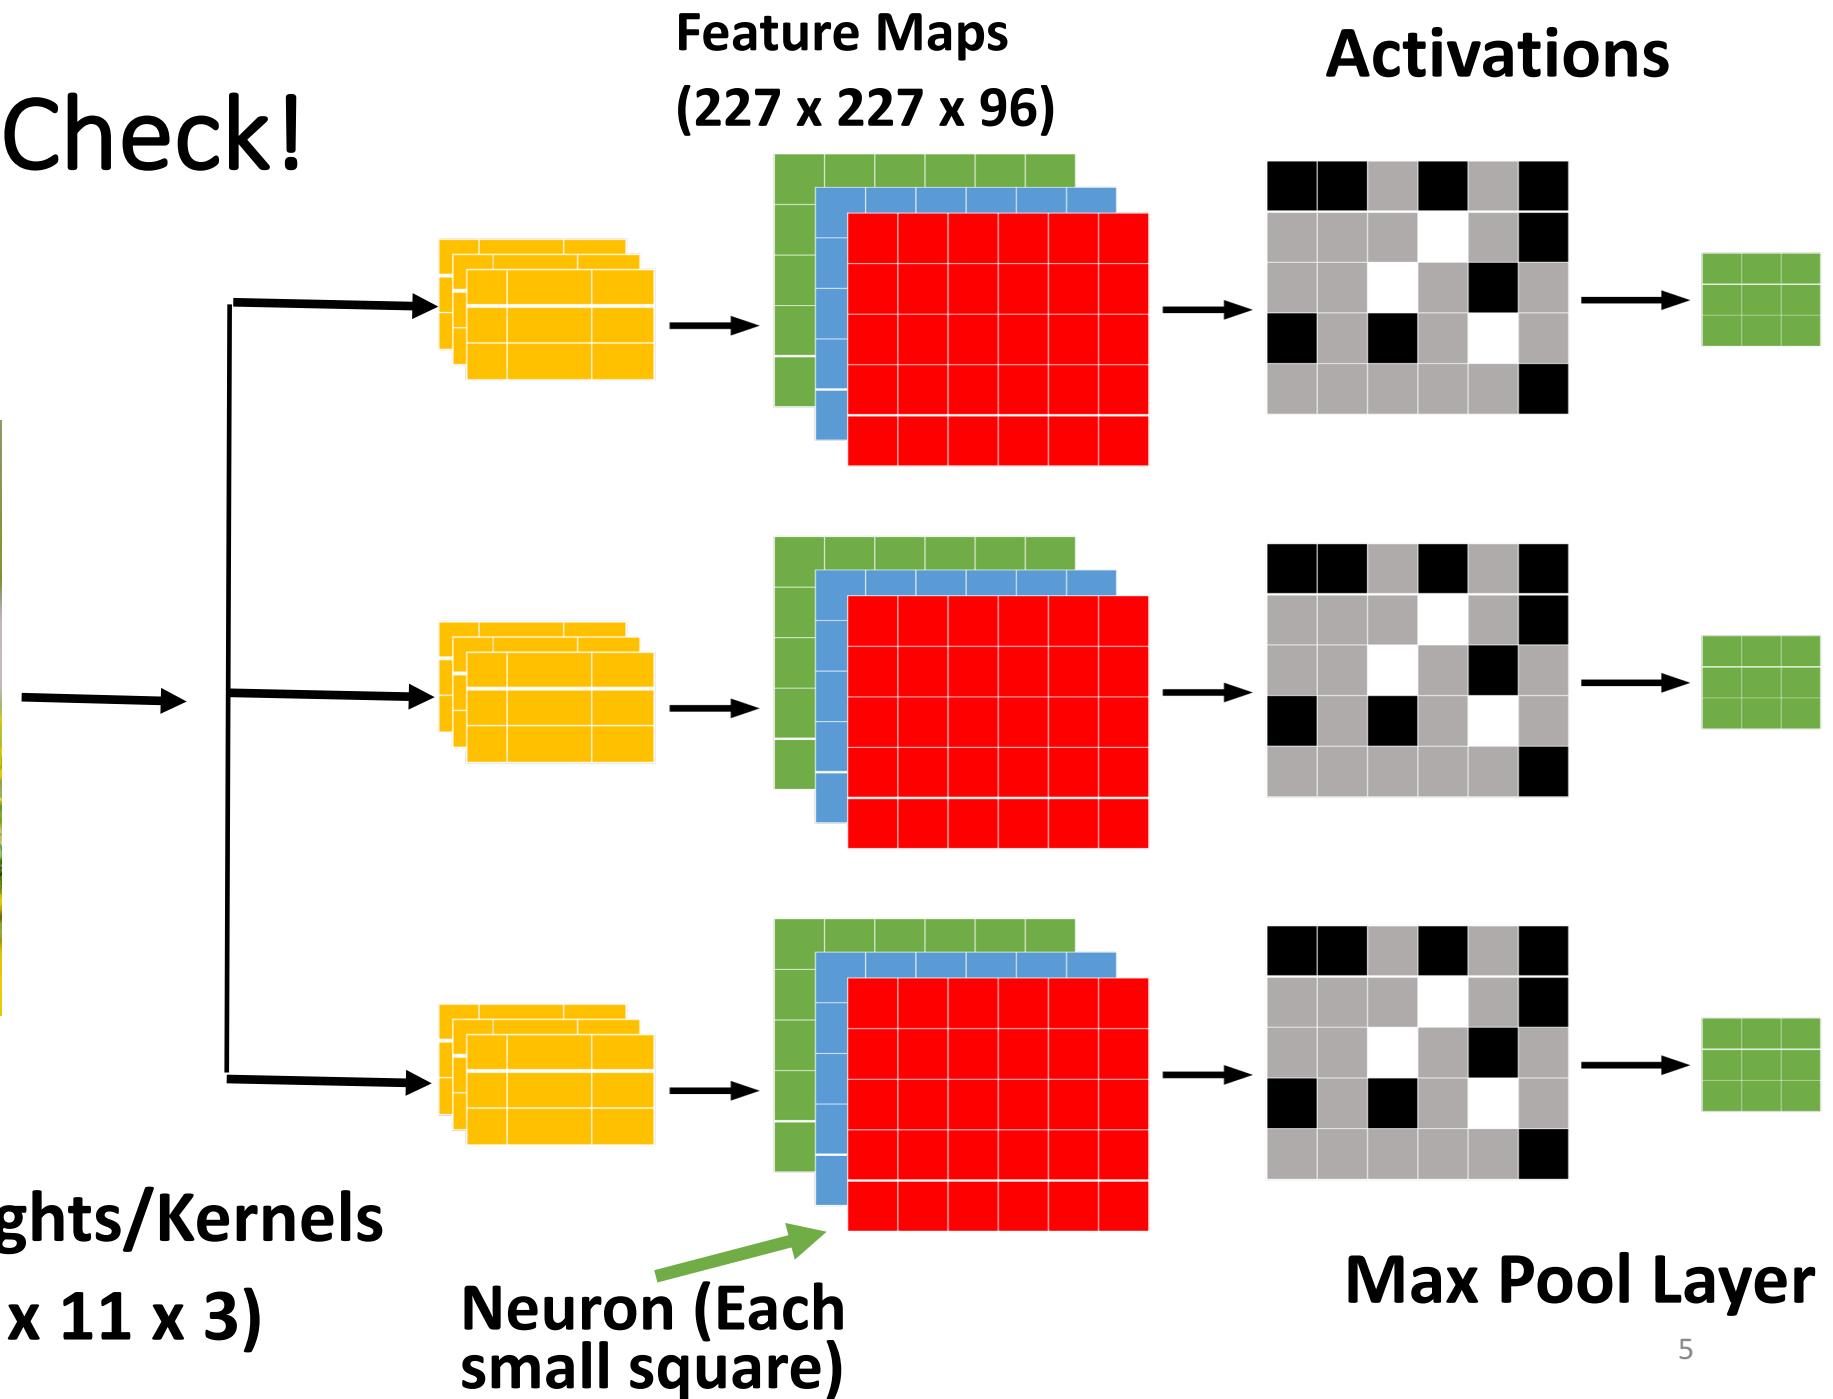
Background Check!

Input Image

($227 \times 227 \times 3$)



Filters/Weights/Kernels
(eg: 96 , $11 \times 11 \times 3$)



Max Pool Layer

What is visualization?

Mapping between a neuron in a layer to the features in the original image.

Backpropagation : How does the loss change with weights?

Visualization : How does the activation of a particular neuron change when we change a part in the image?

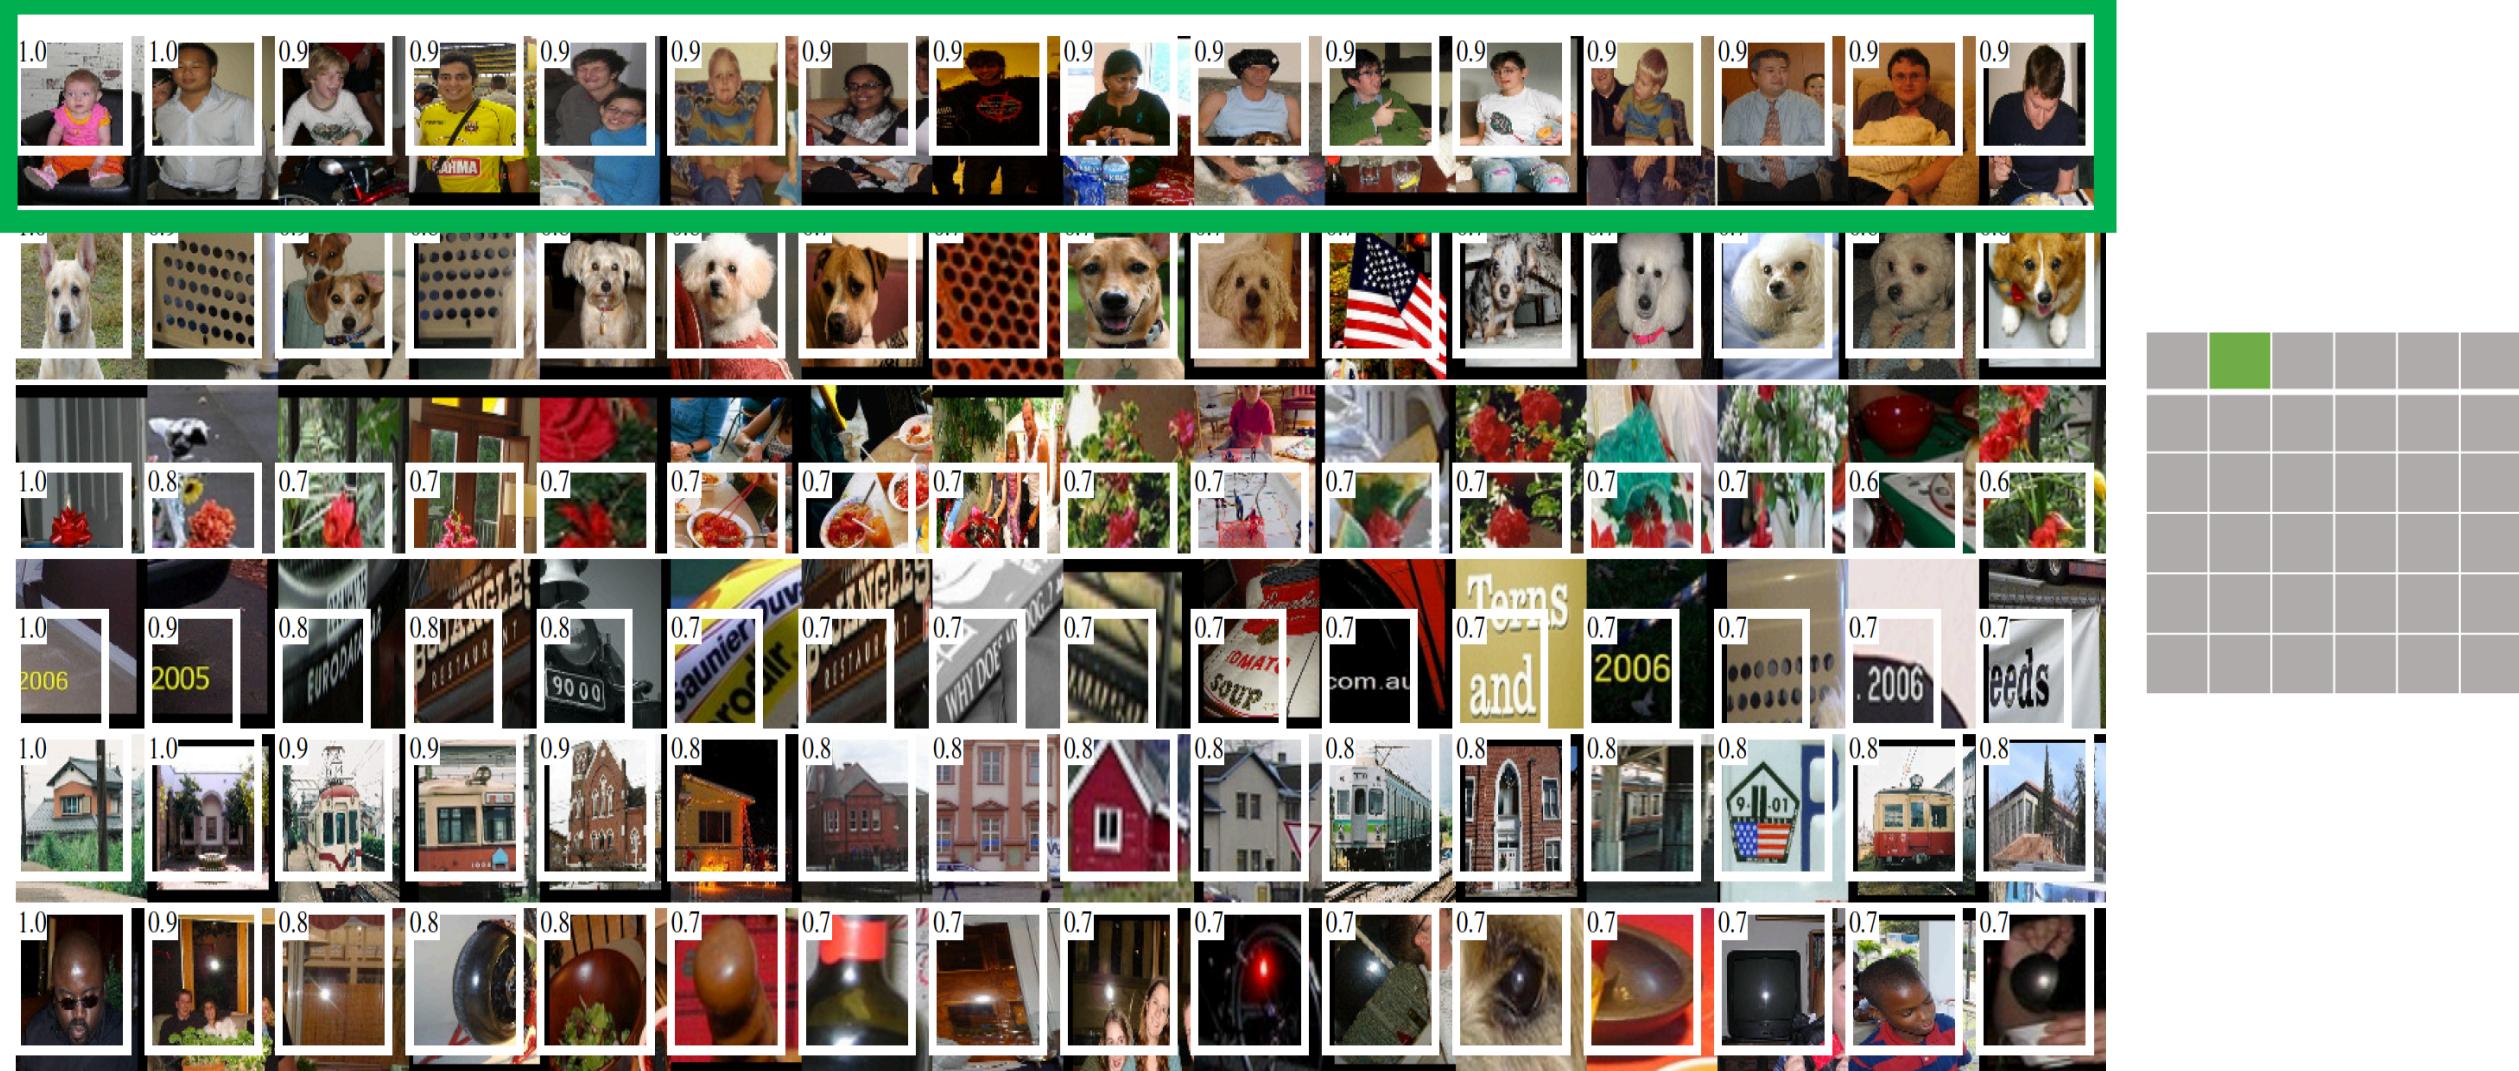
Why visualization?

- Understand how and why neural networks work
- Observe the evolution of features during training
- Aid the development of better models (rather than just trial-and-error)
- Diagnose potential problems with the model

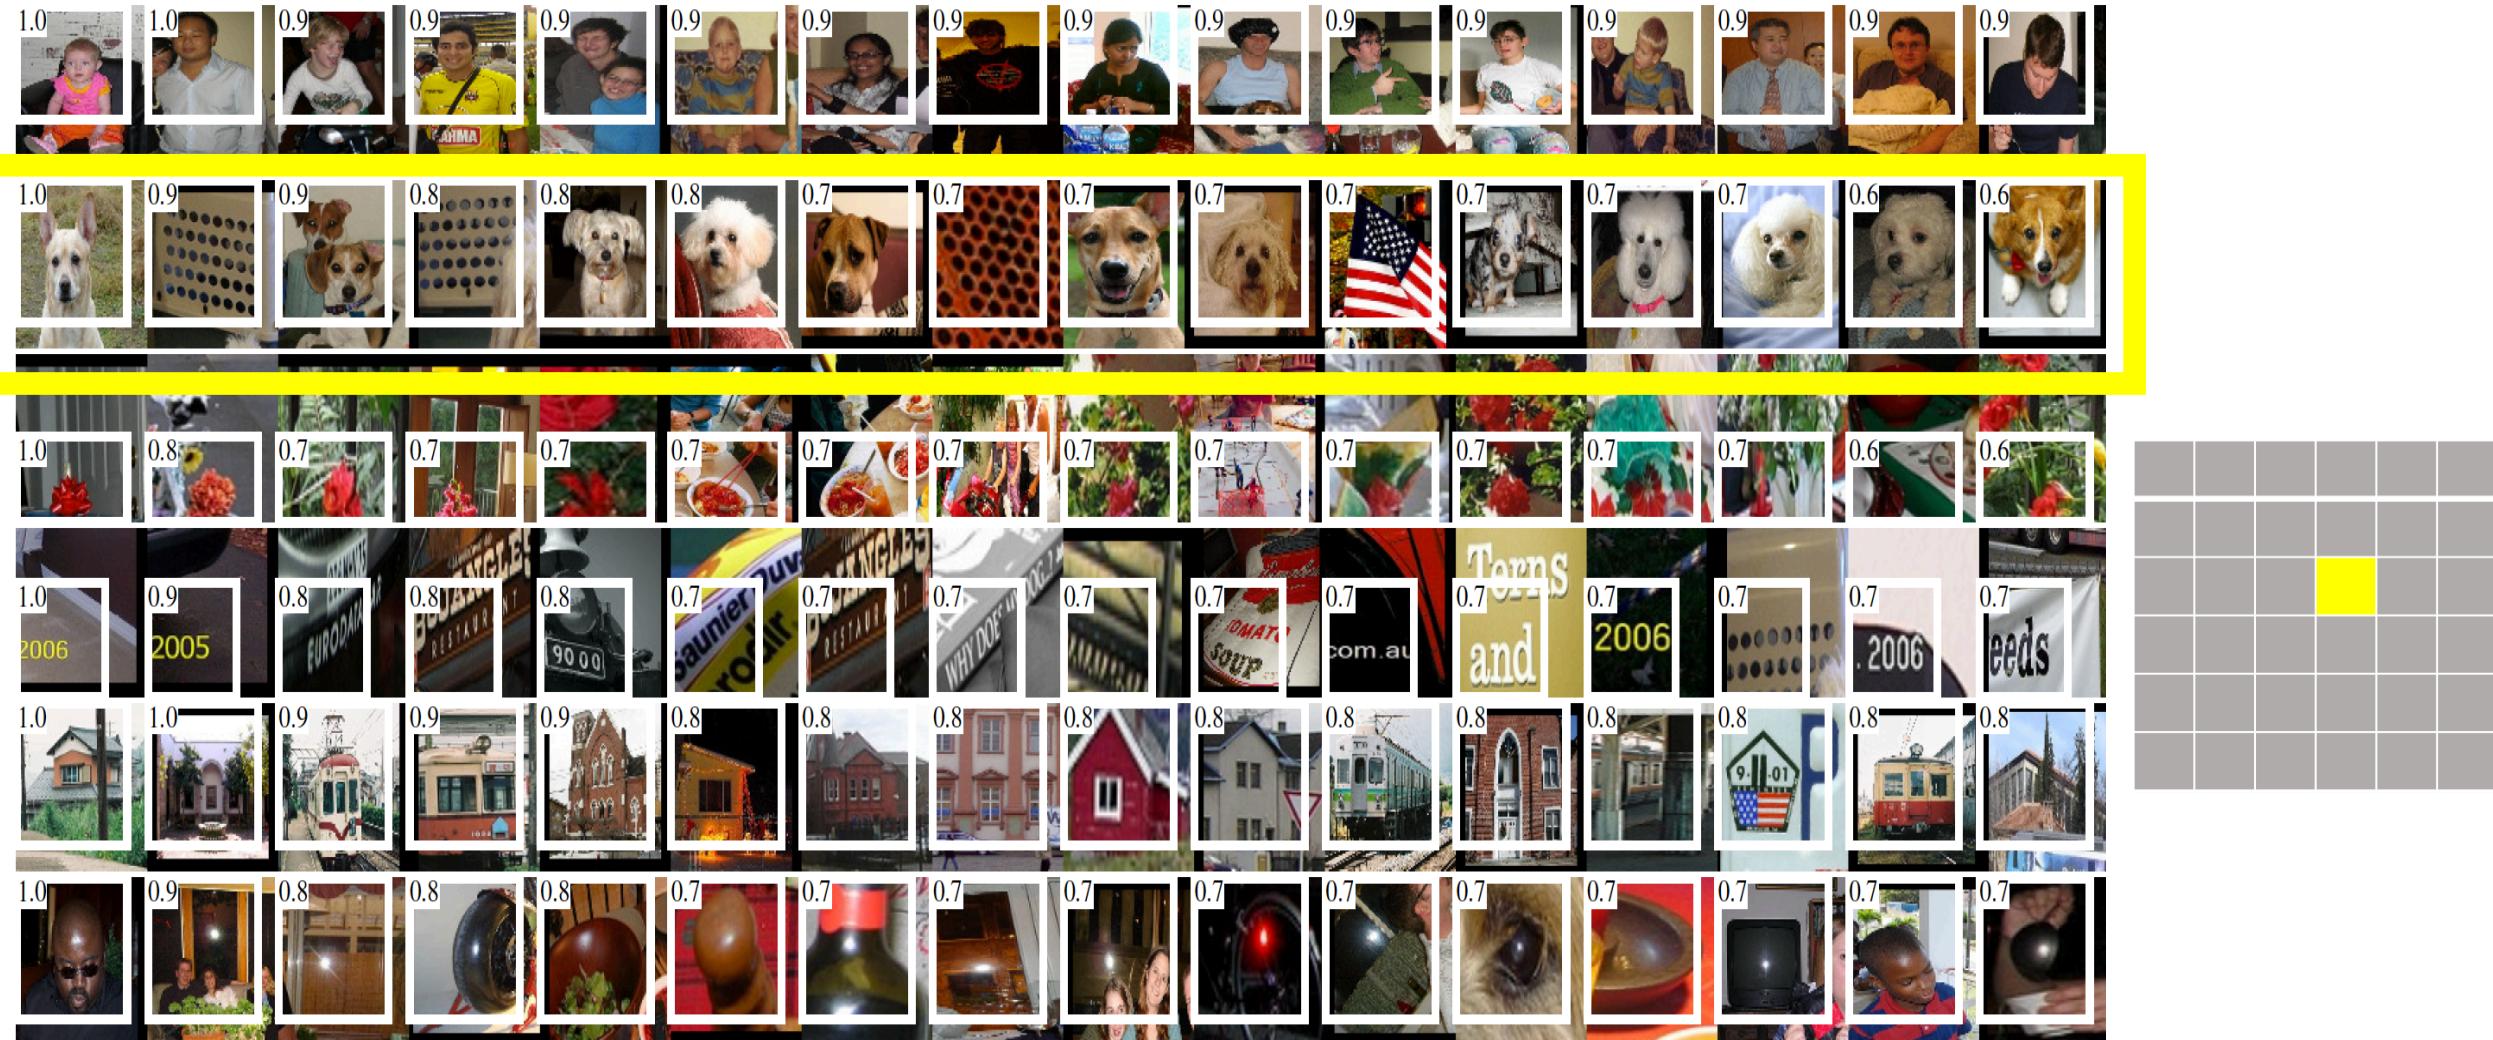
Outline - Visualization

- What is Visualization?
- Visualize patches that maximally activate neurons
- Visualize the weights
- Gradient based approaches
- Optimization based approach

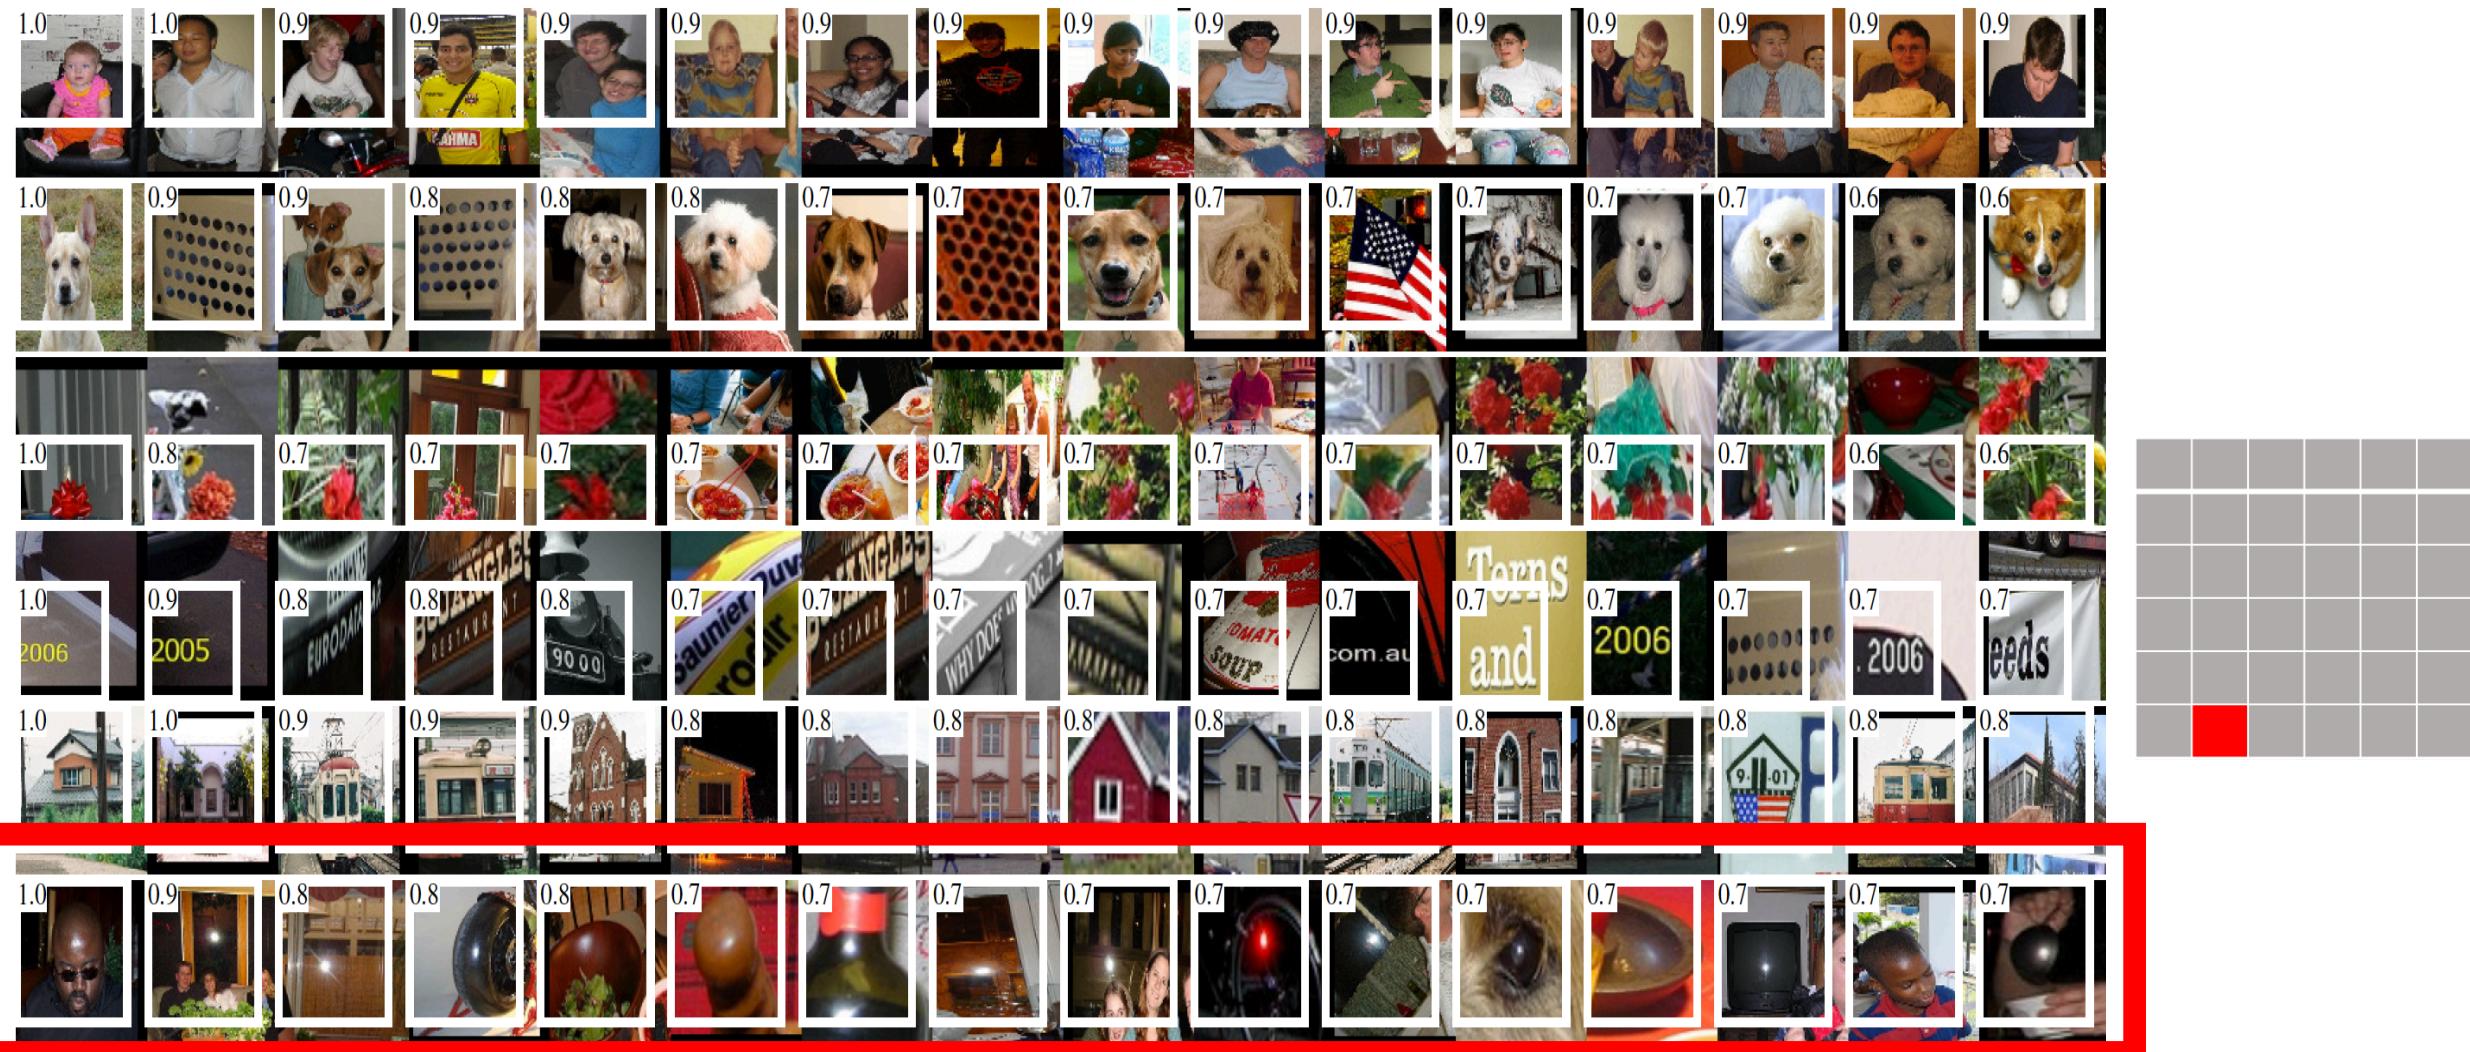
Visualize patches that maximally activate neurons



Visualize patches that maximally activate neurons



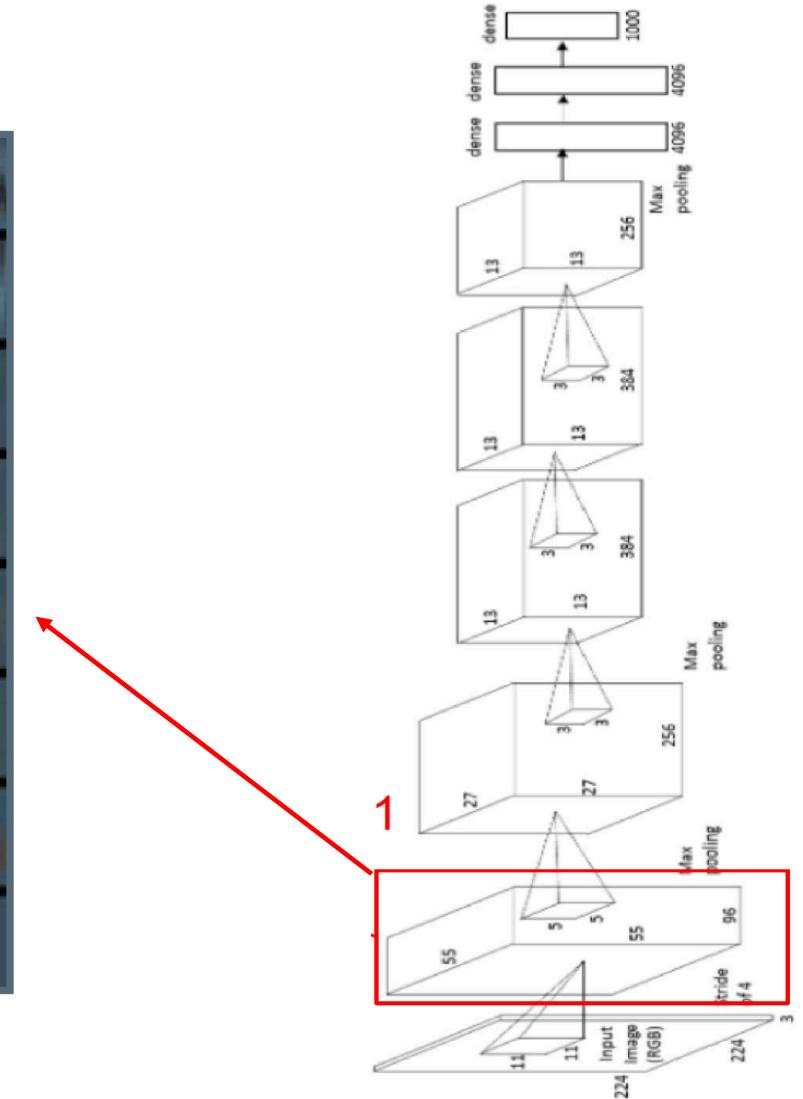
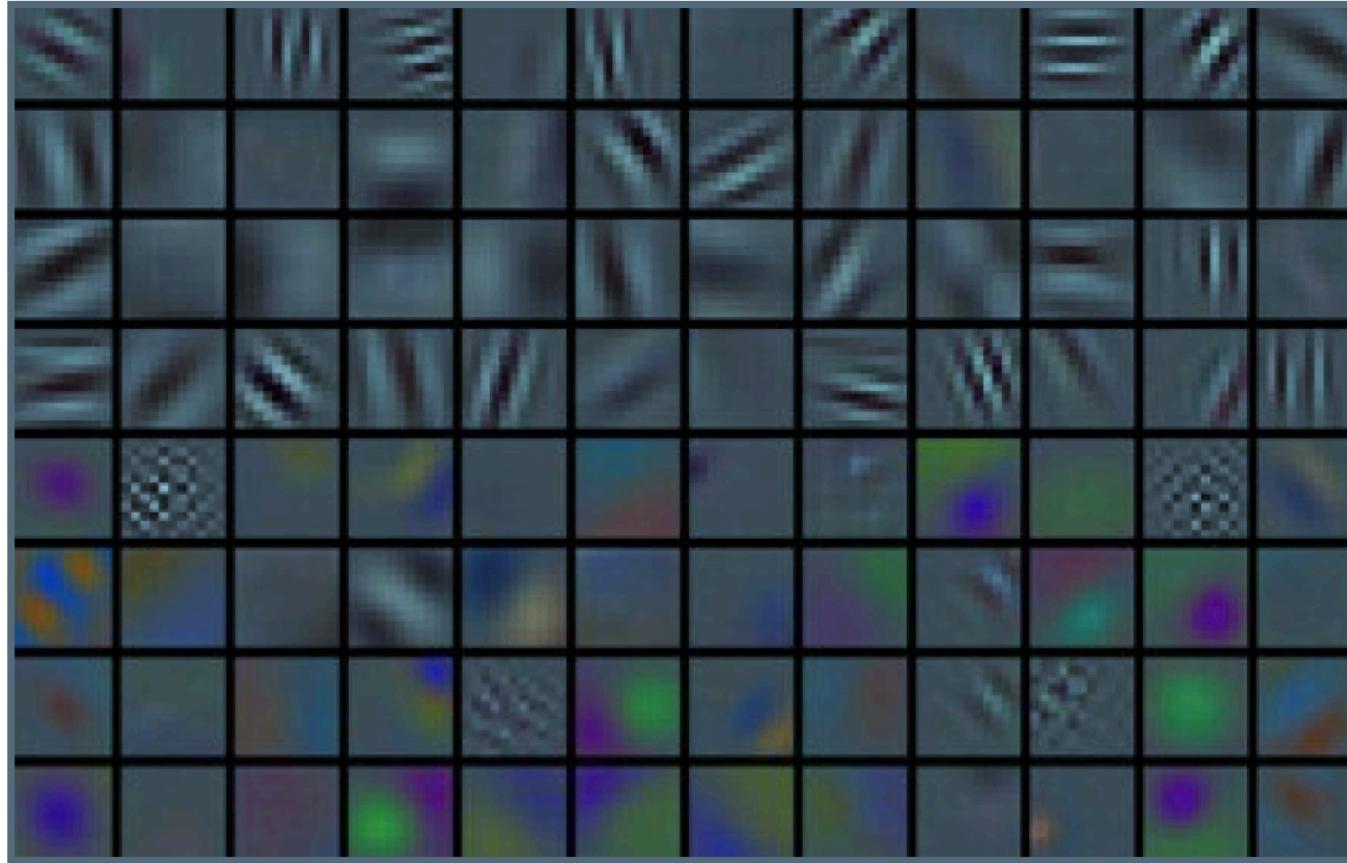
Visualize patches that maximally activate neurons



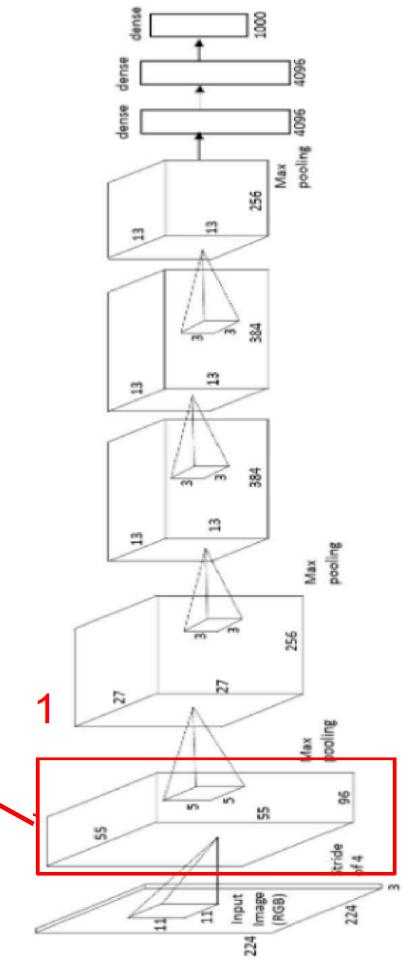
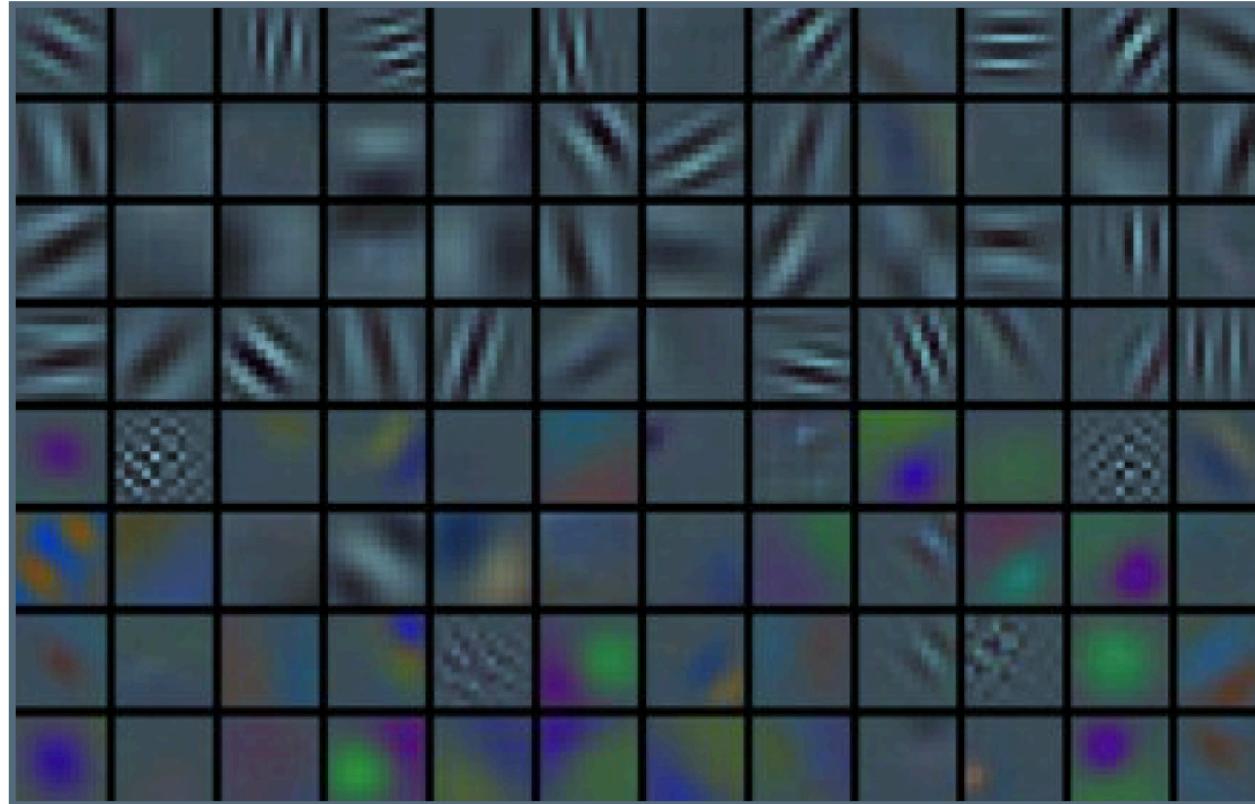
Outline - Visualization

- What is Visualization?
- Visualize patches that maximally activate neurons
- **Visualize the weights**
- Gradient based approaches
- Optimization based approach

Visualize the weights



Visualize the weights



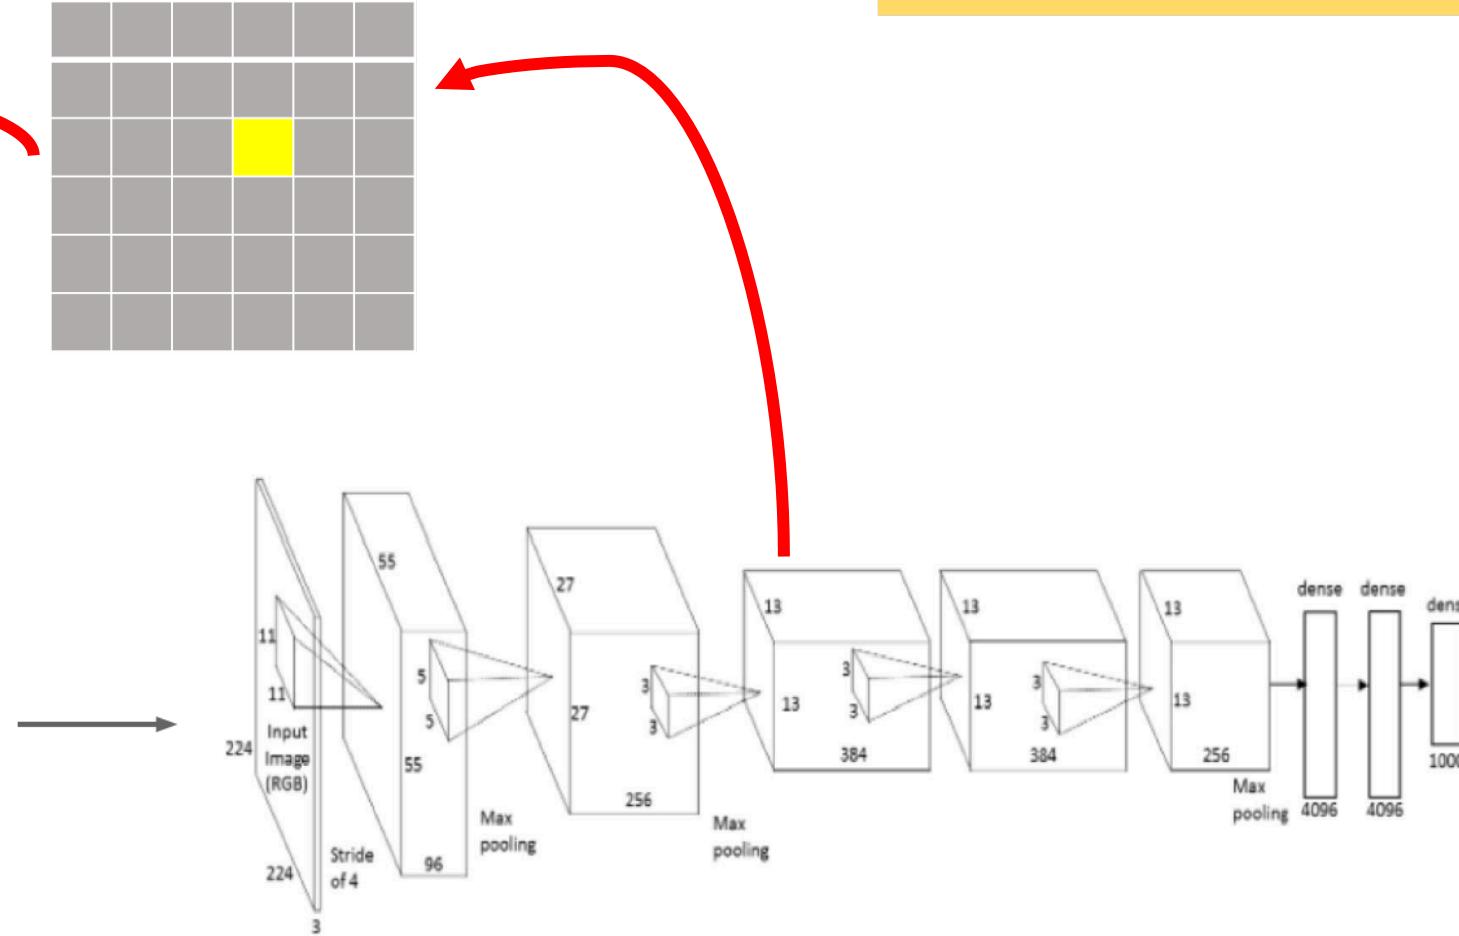
Only possible for the first layer 😞

Outline - Visualization

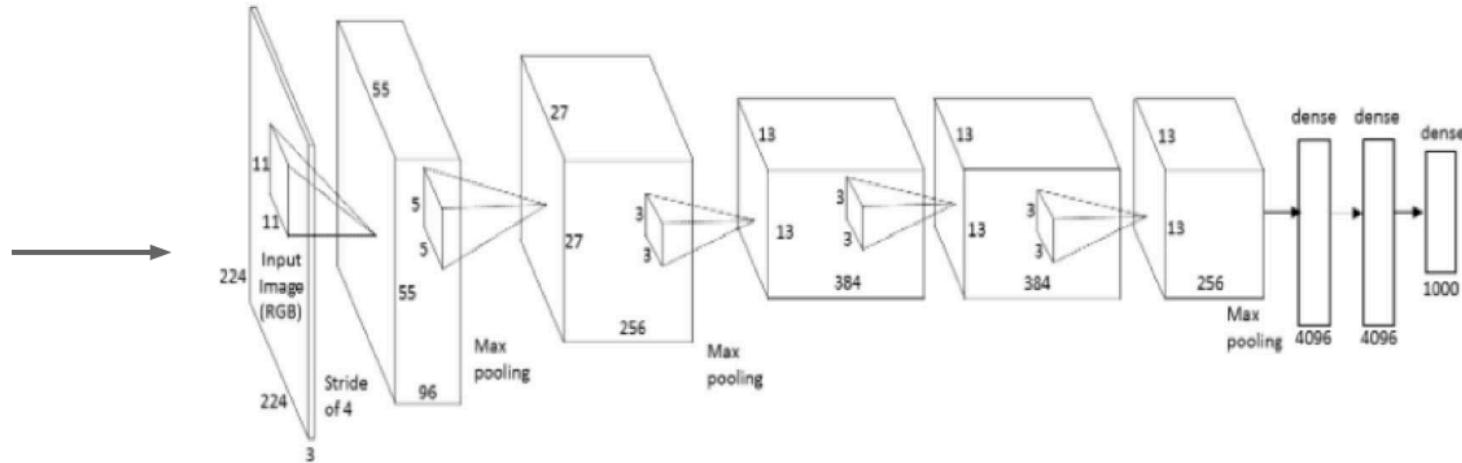
- What is Visualization?
- Visualize patches that maximally activate neurons
- Visualize the weights
- **Gradient based approaches**
- Optimization based approach

Gradient based approaches

Q : How can we compute the gradient of an arbitrary neuron w.r.t the image?



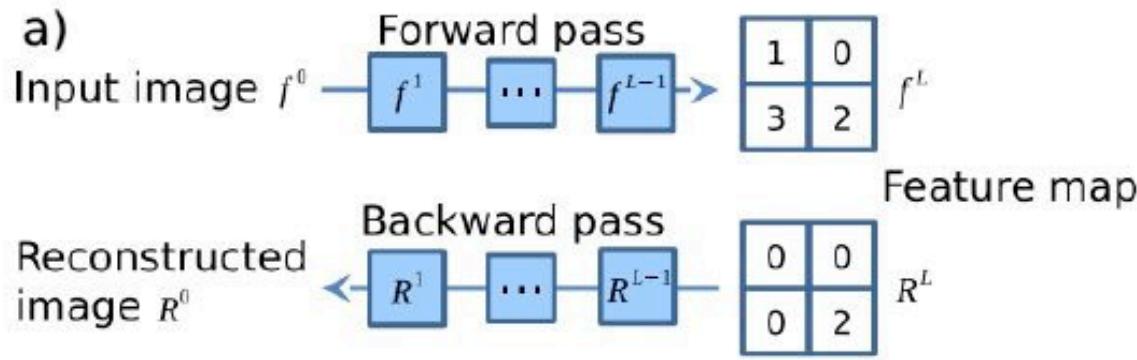
Gradient based approaches



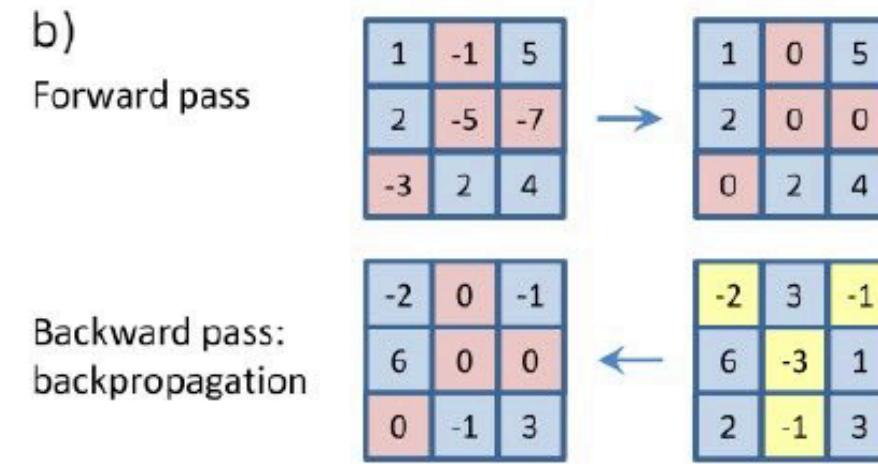
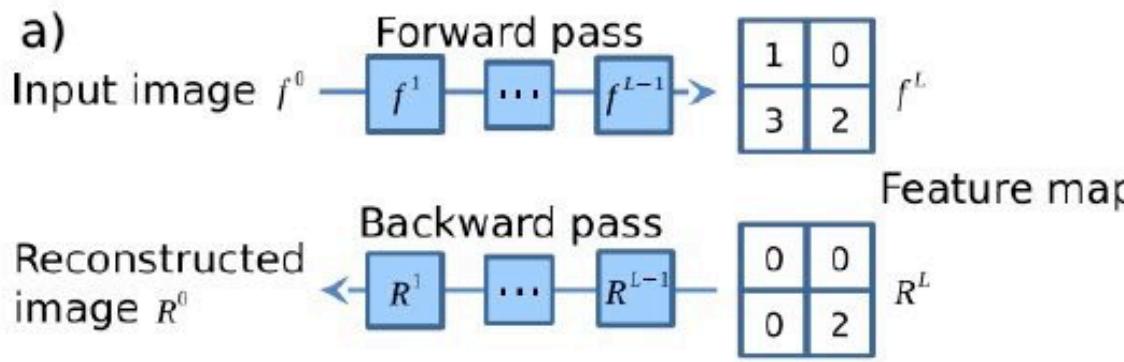
1. Input the image into the net
2. Pick a layer, set the **gradient** there to be all **0** except for one **1** for some neuron of interest
3. “Map it” back to the image



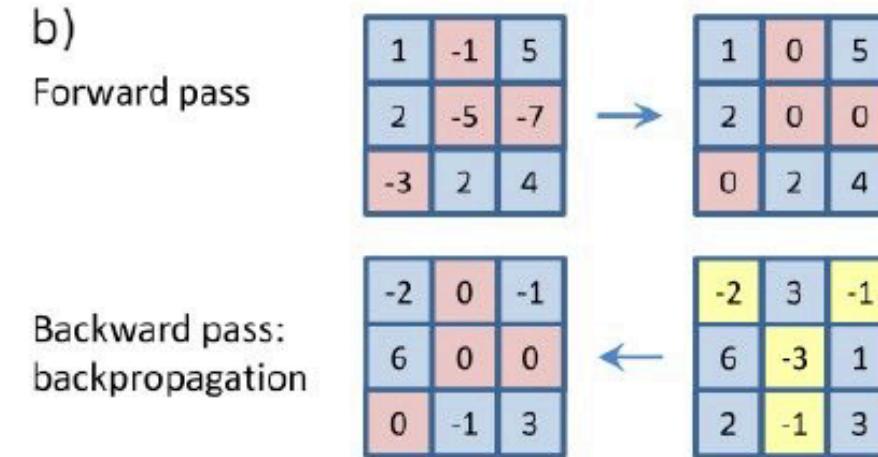
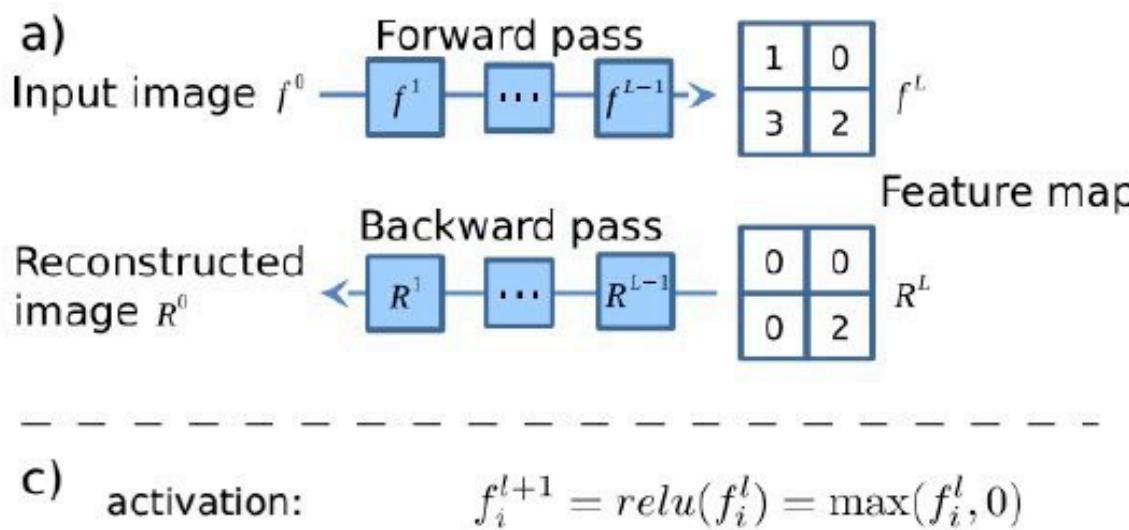
Gradient based approaches - “Map back”



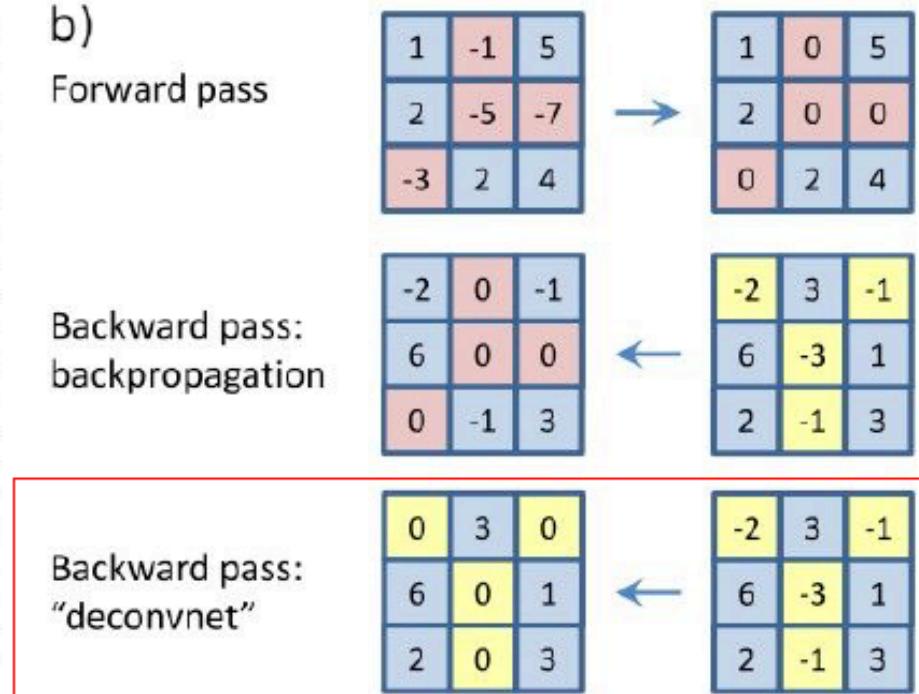
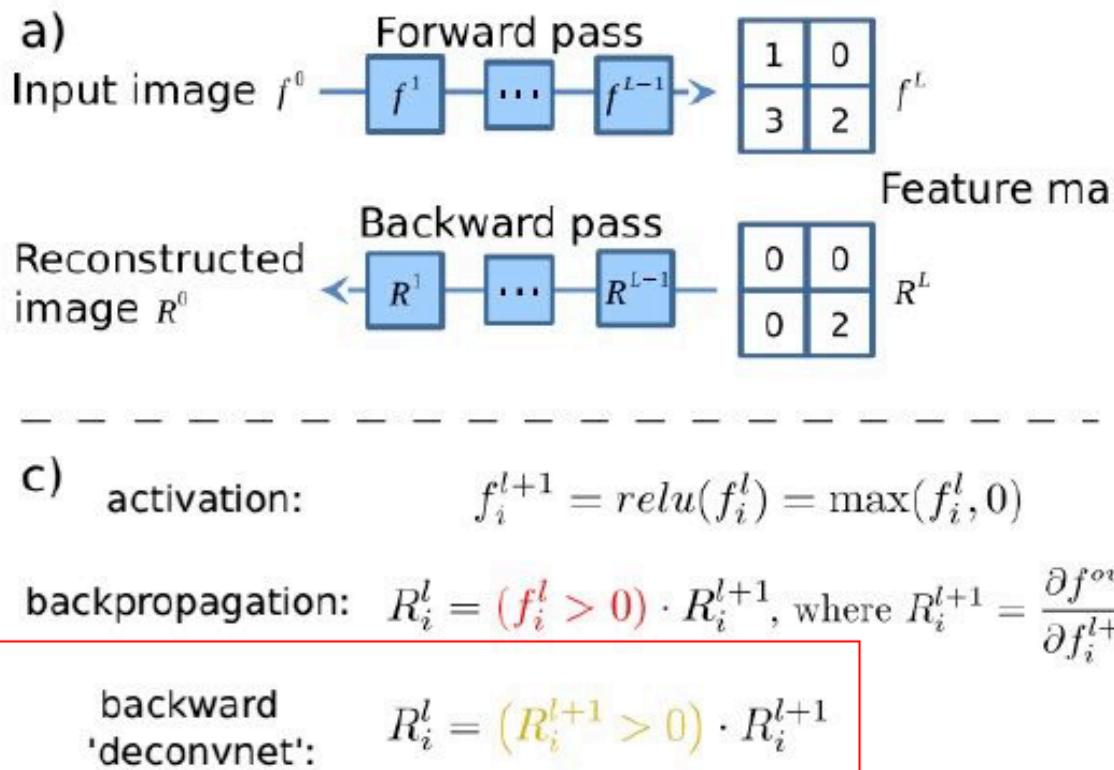
Gradient based approaches - “Map back”



Gradient based approaches - “Map back”



Gradient based approaches - “Map back”



Deconvnet !

Visualizing the neurons along
the way to the top

Choose a target neuron



Input the images 1 by 1



Select the top 9 images that have the highest
activation for that neuron

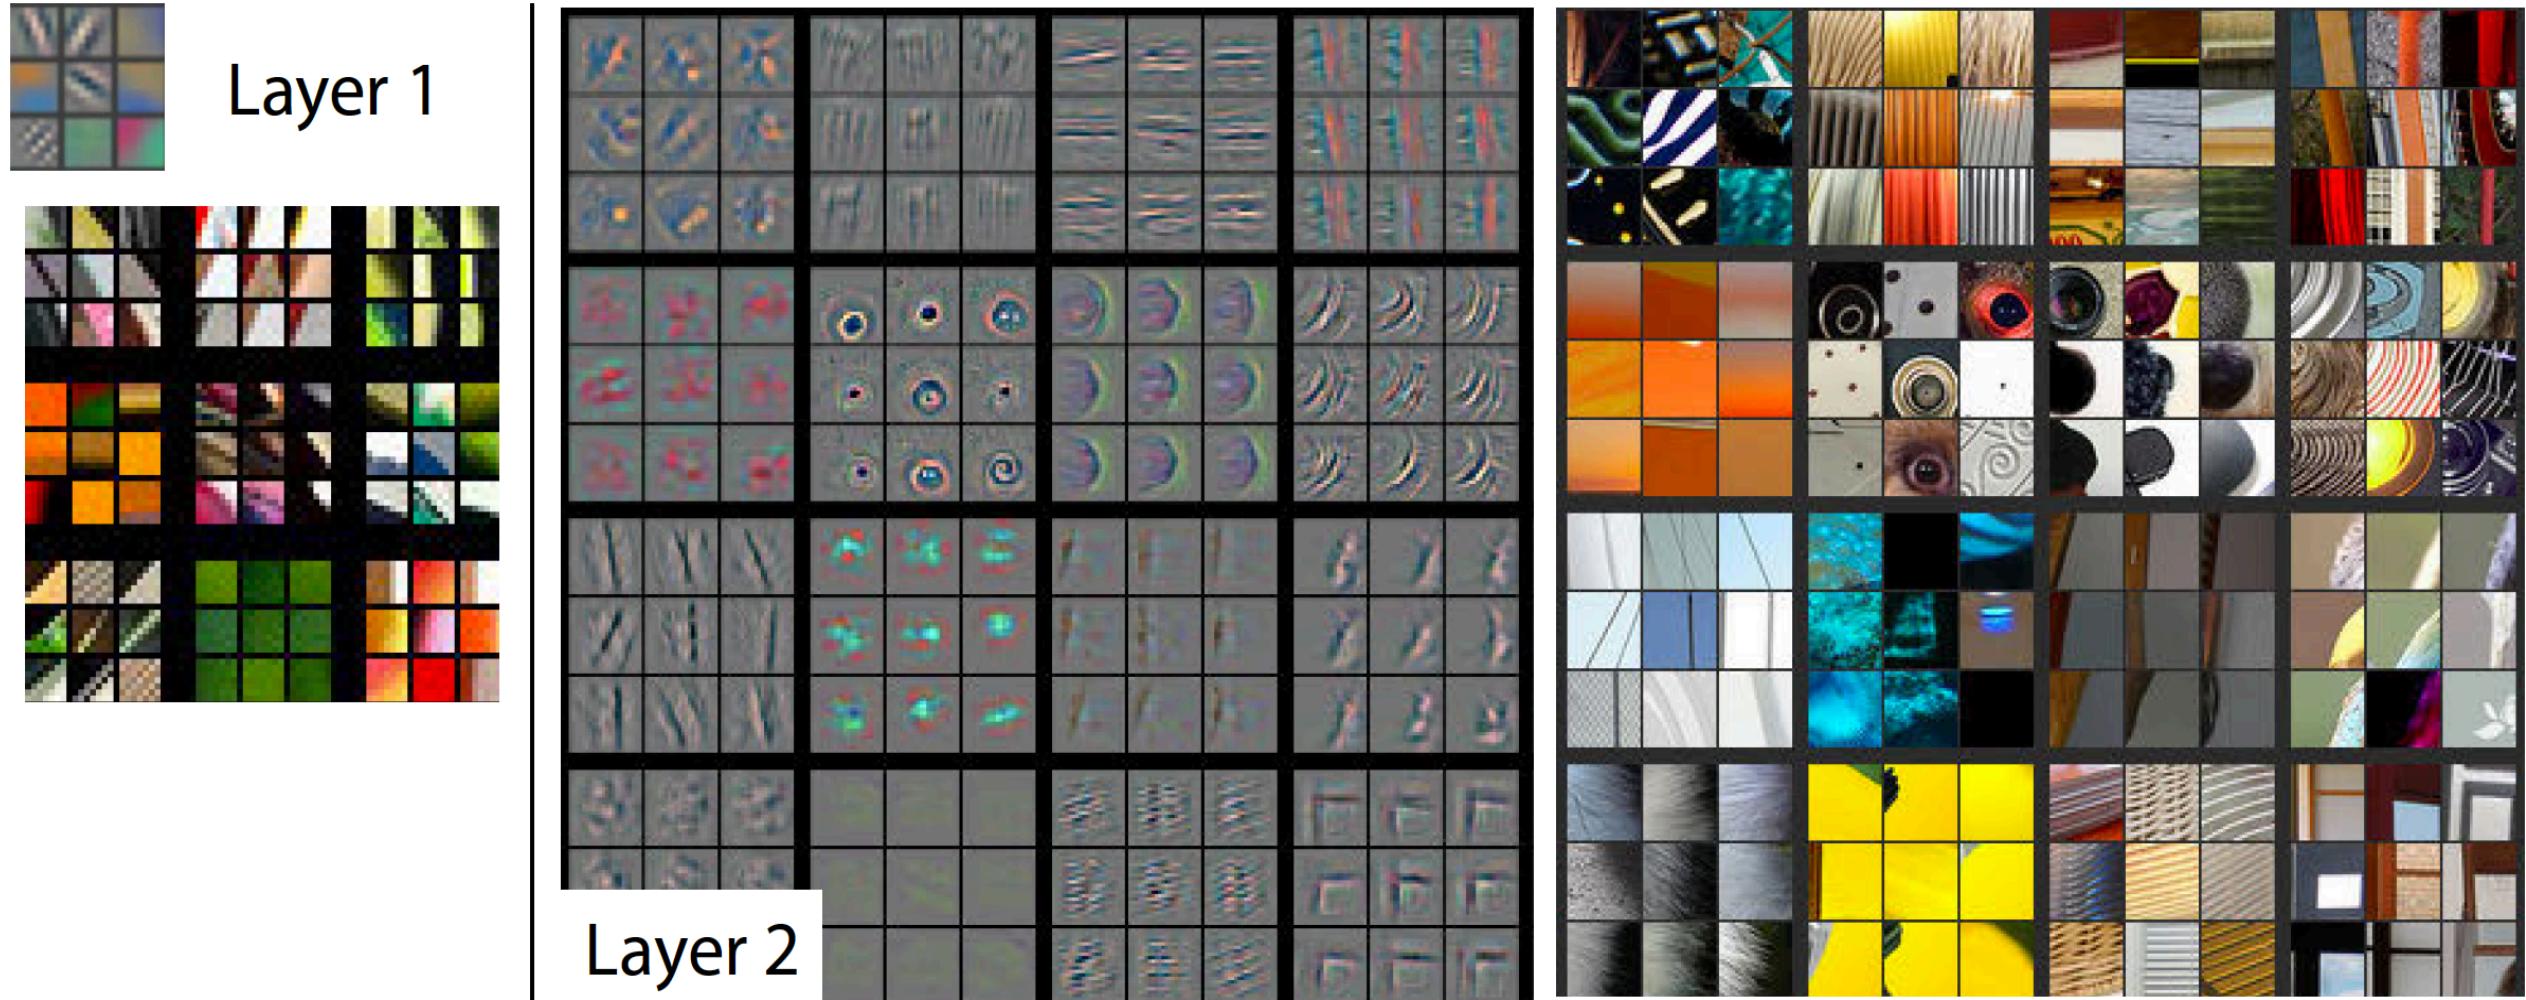


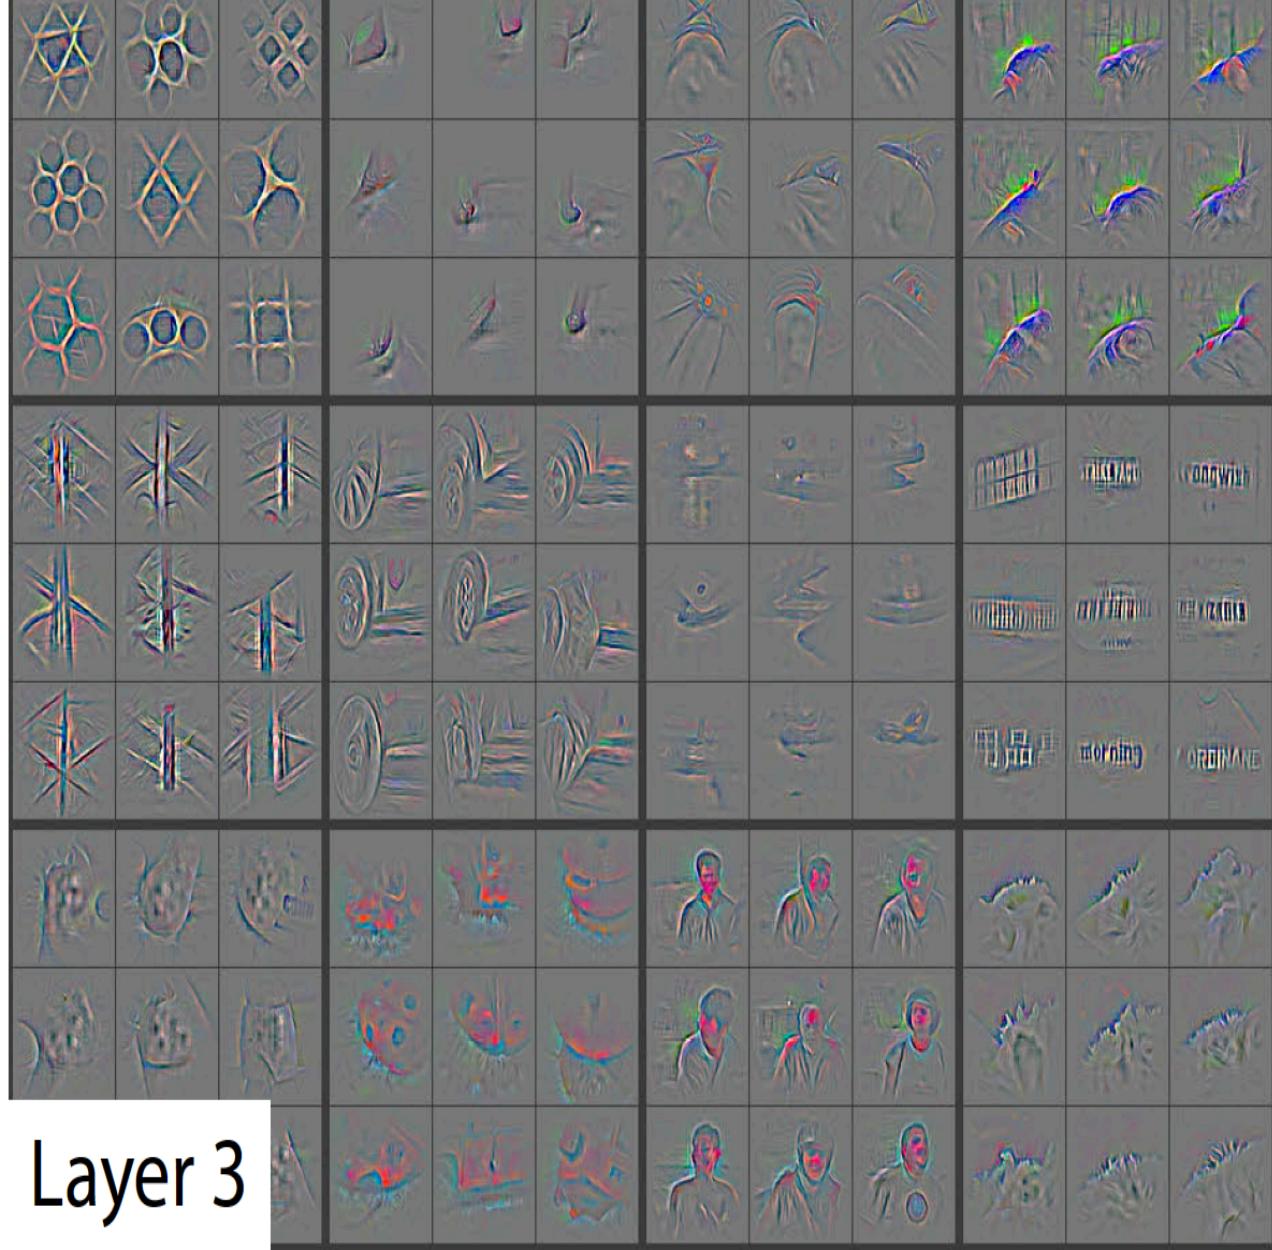
Cluster those images together



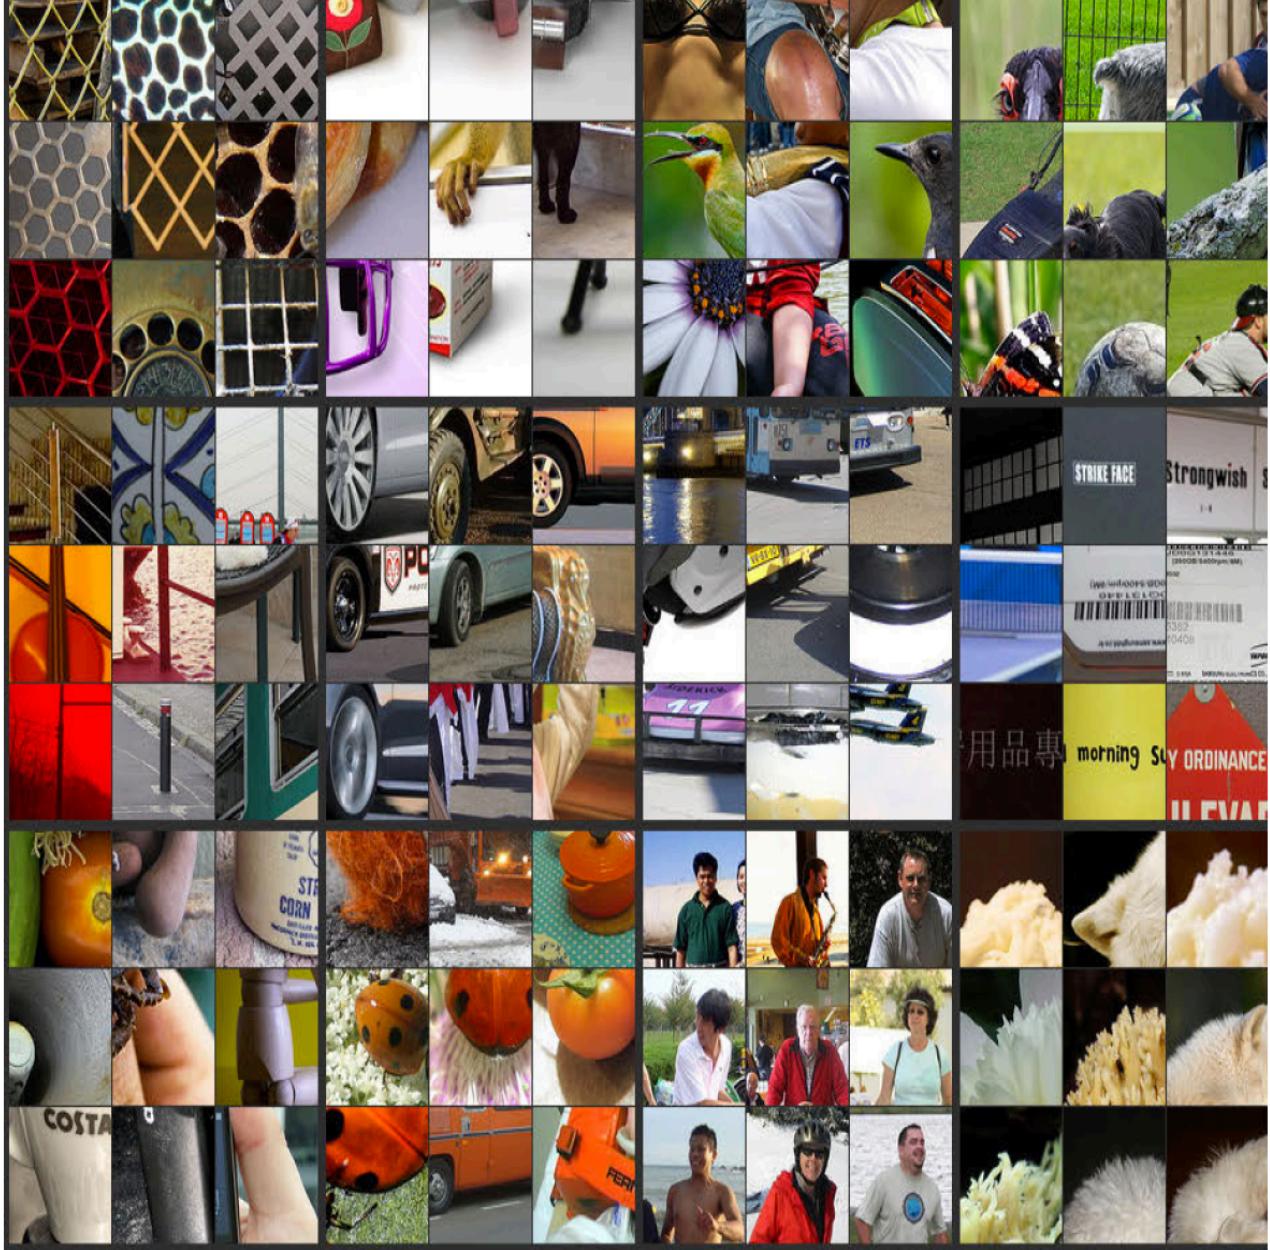
Map back from that neuron and create a “back-
pass map”

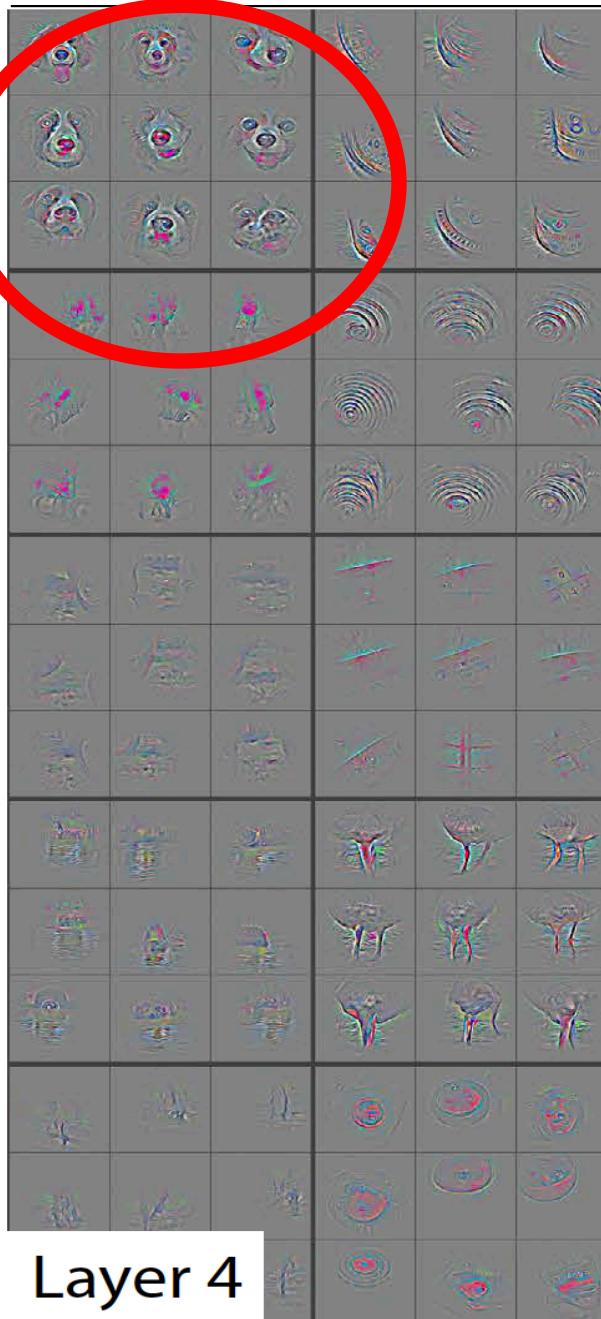
Visualizing the neurons along the way to the top



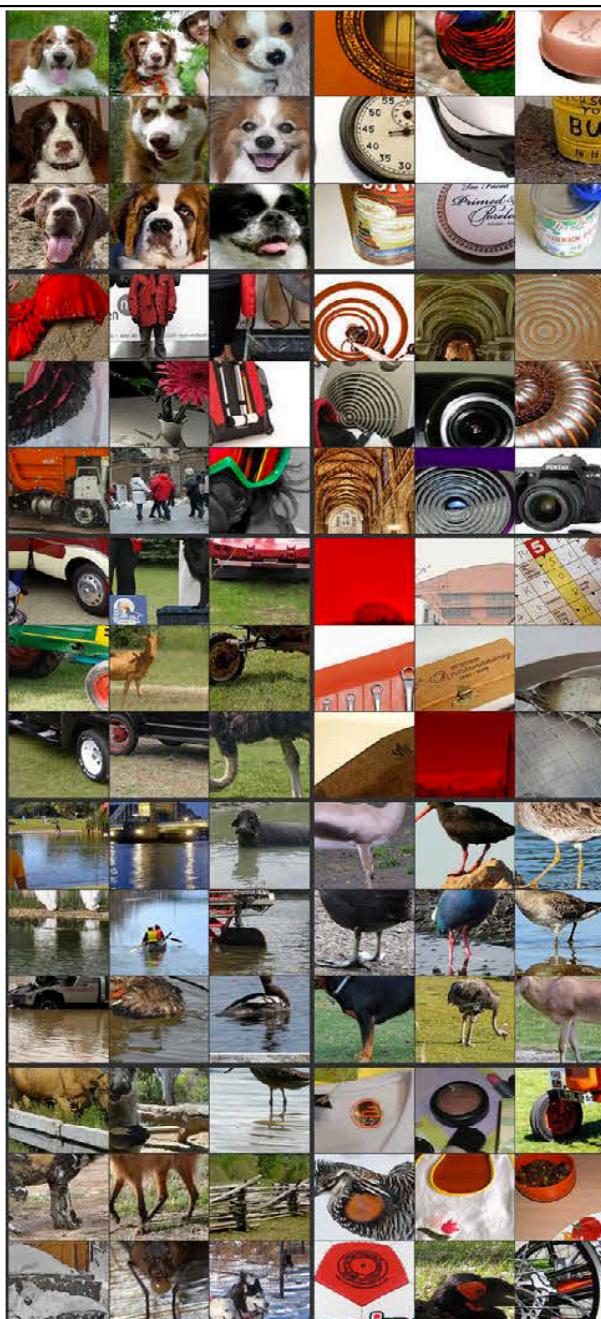


Layer 3

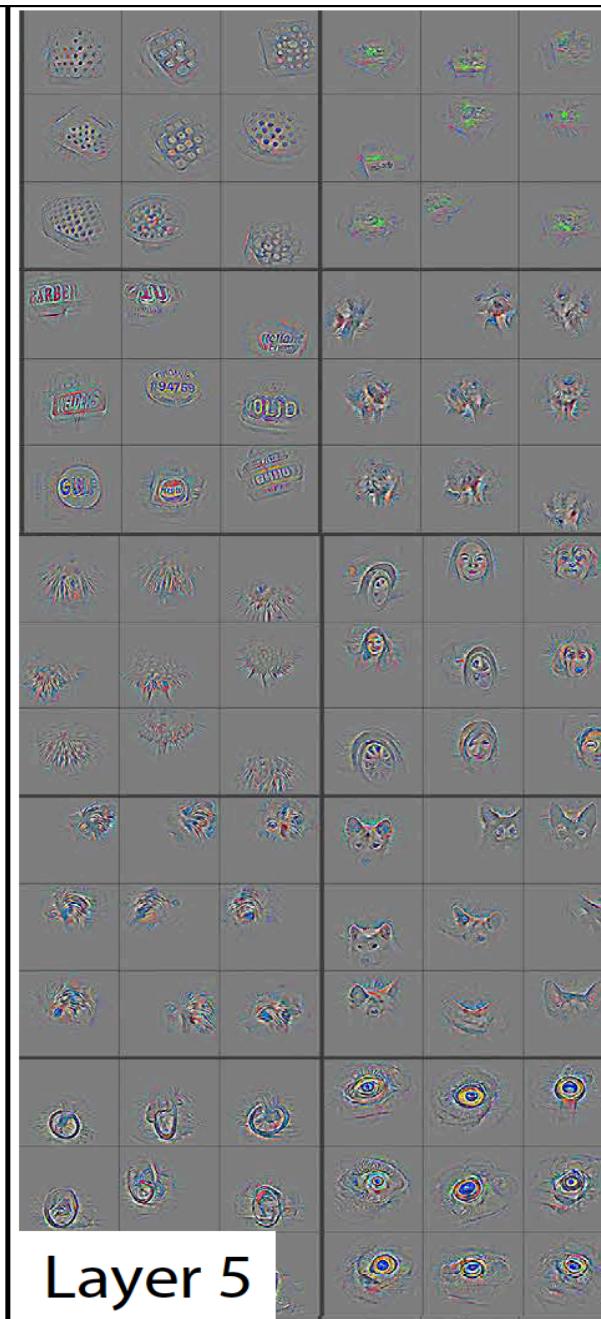




Layer 4

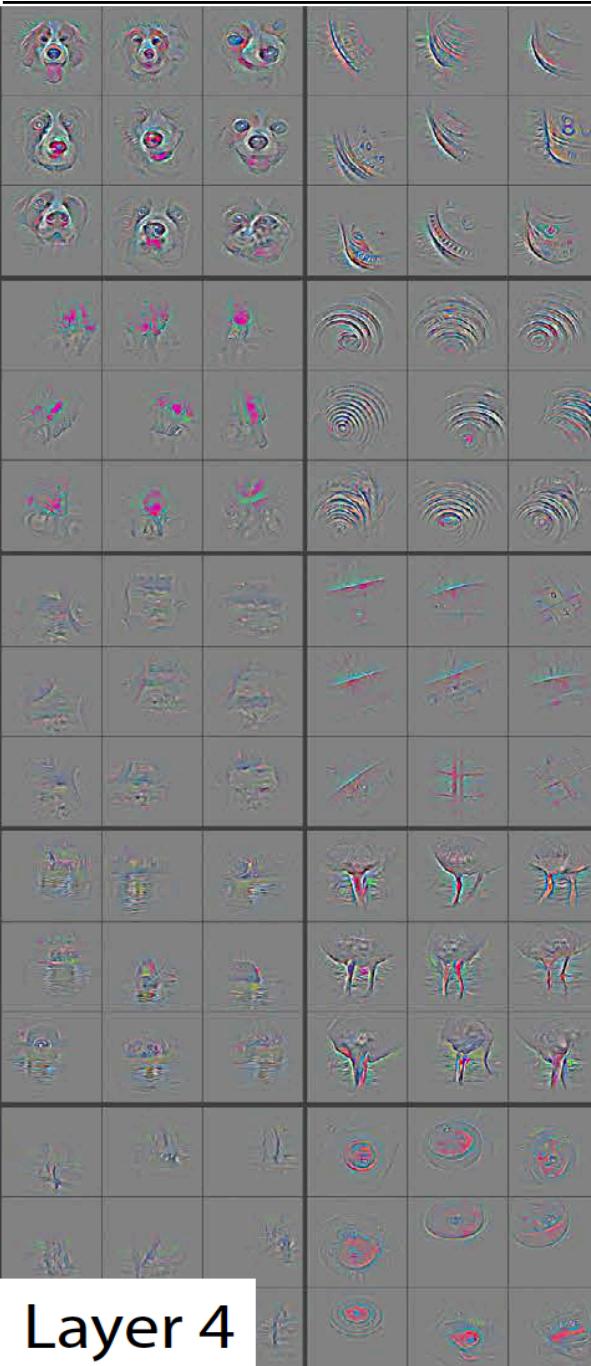


Matthew D. Zeiler, Rob Fergus [Visualizing and Understanding Convolutional Networks](#), ECCV 2014

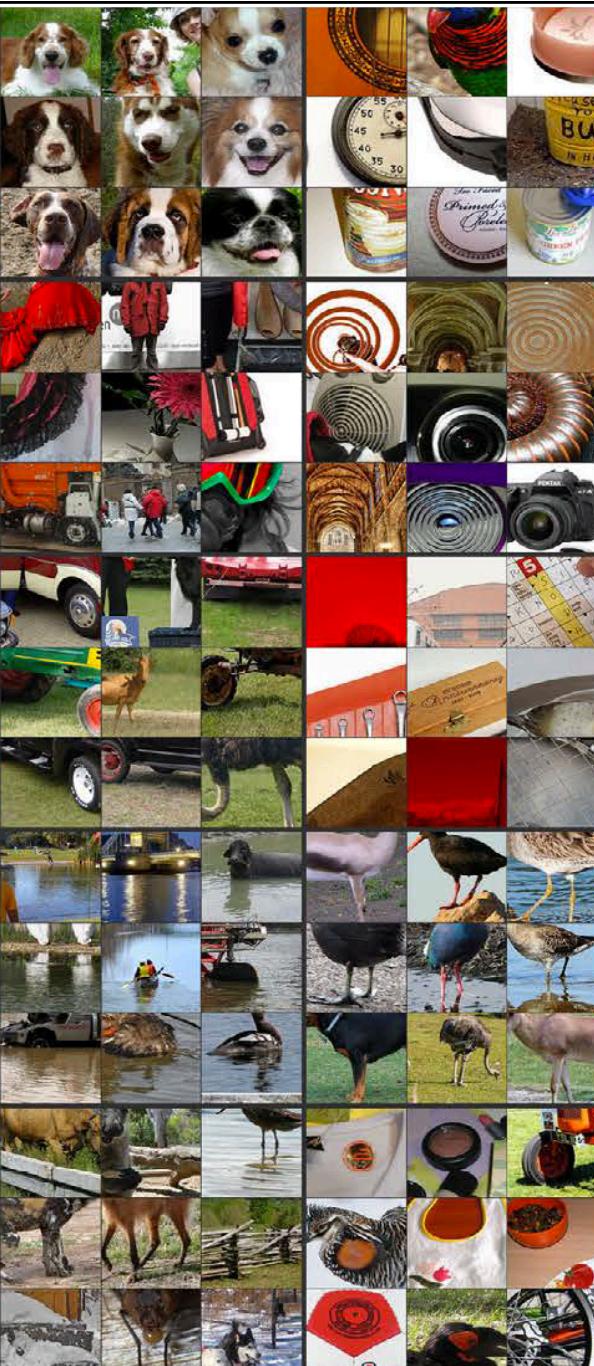


Layer 5

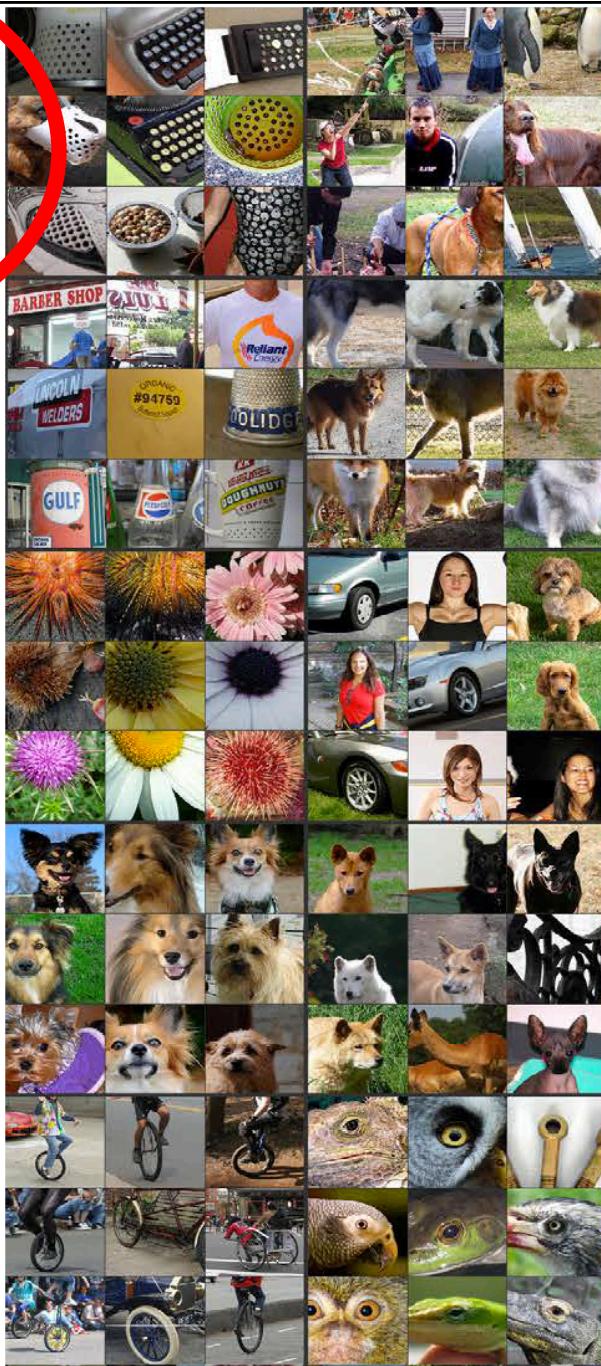
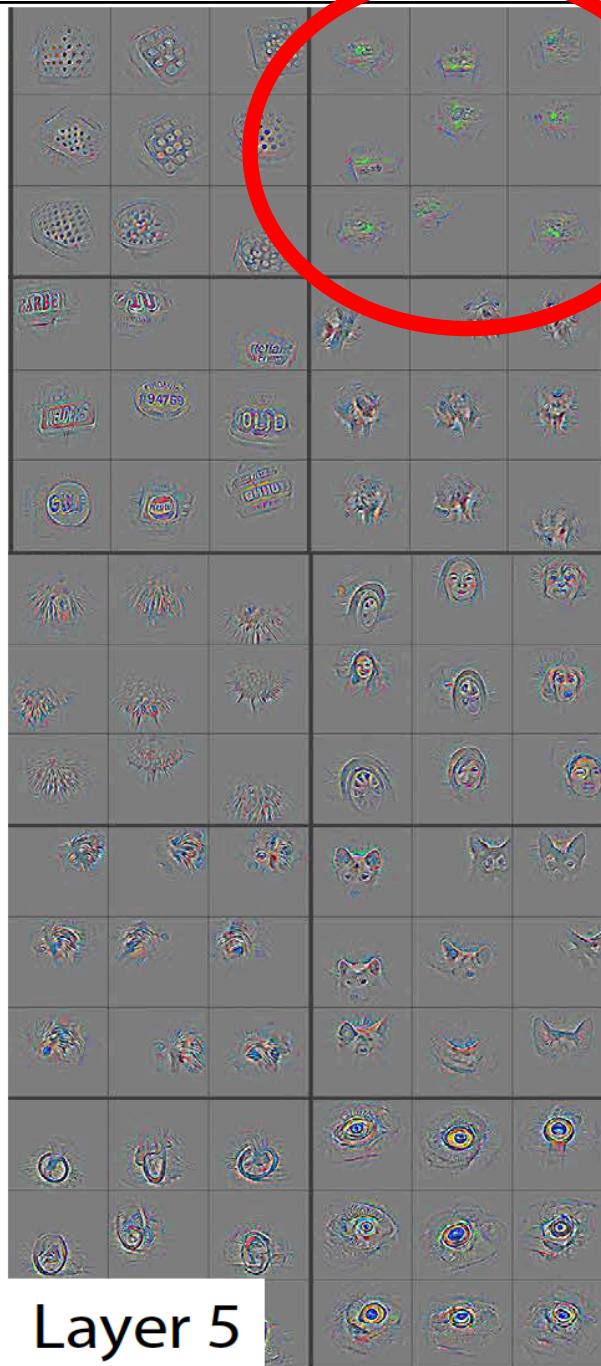




Layer 4



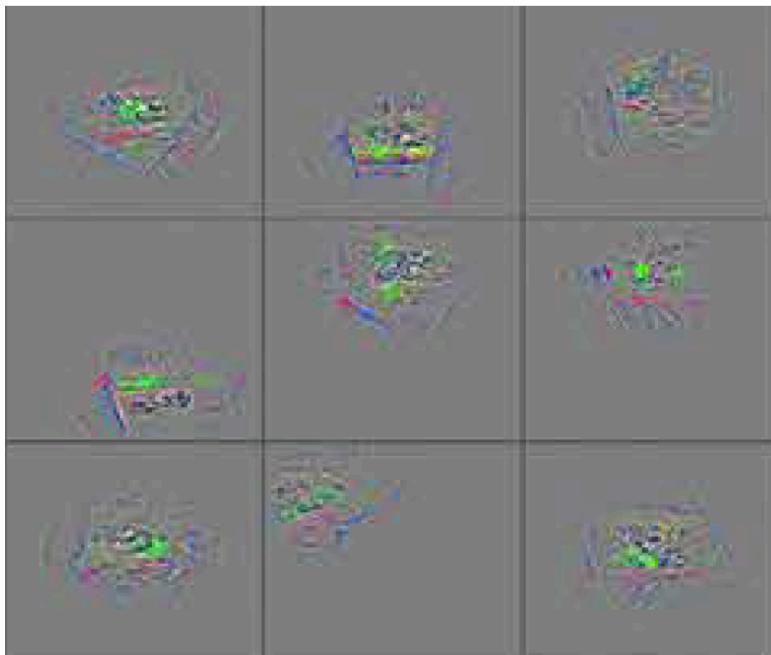
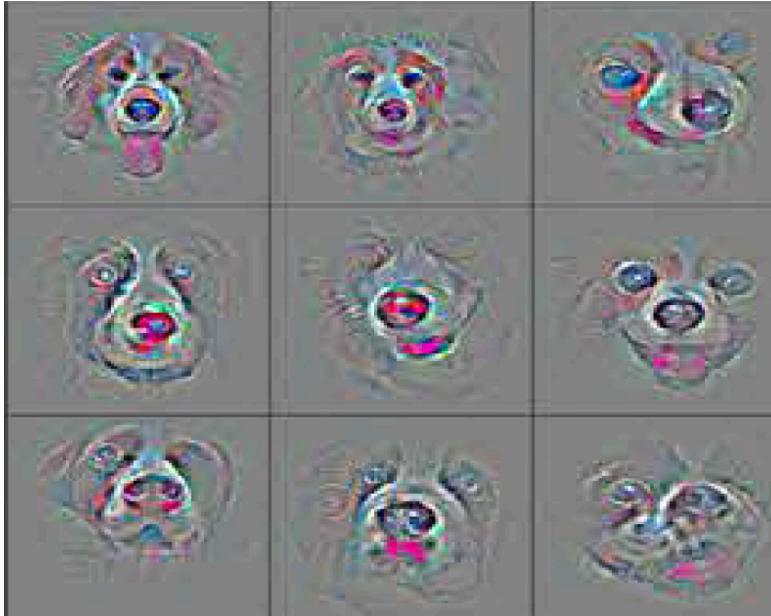
Layer 5

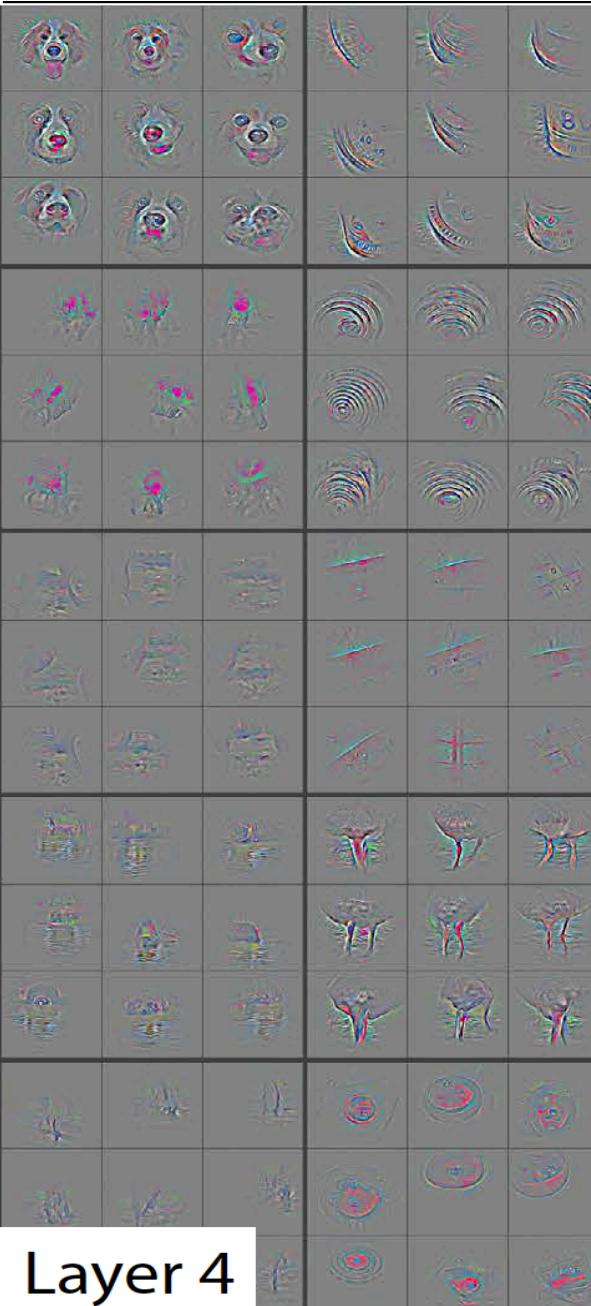




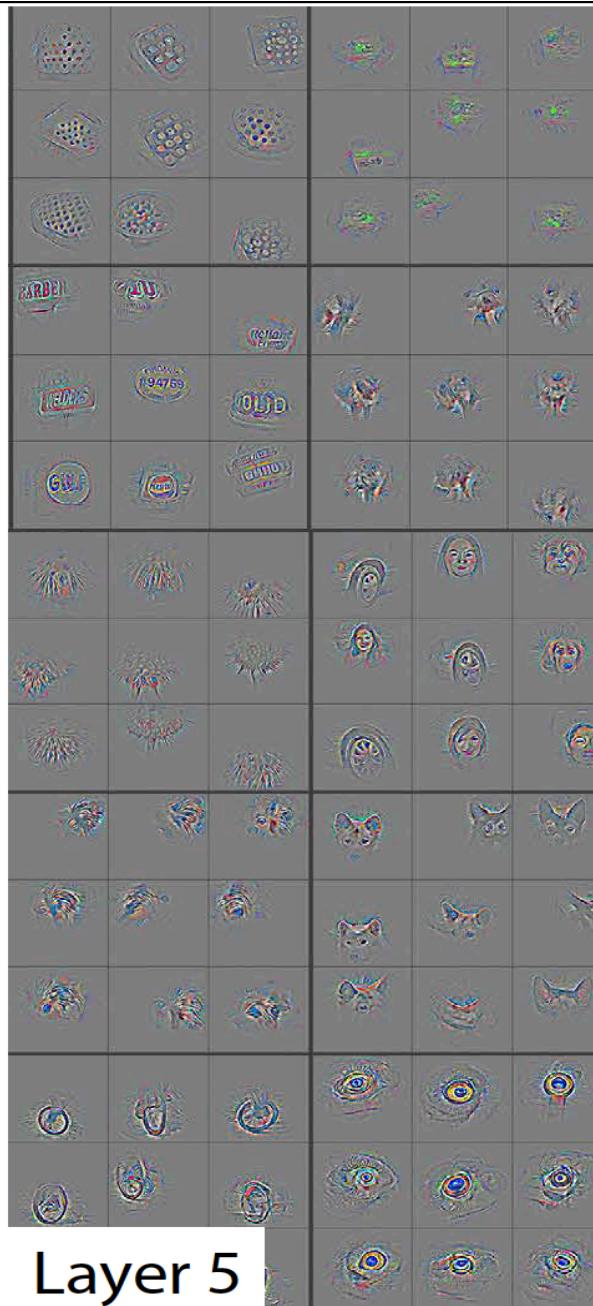
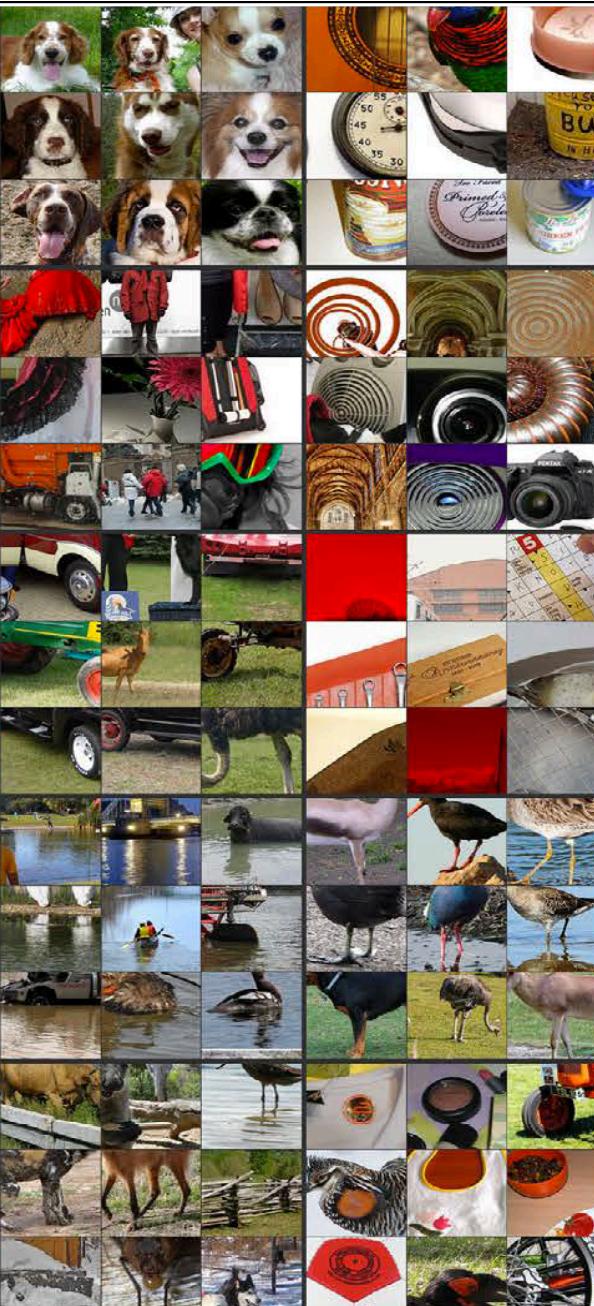
What features are being captured
from these pictures?







Layer 4



Layer 5

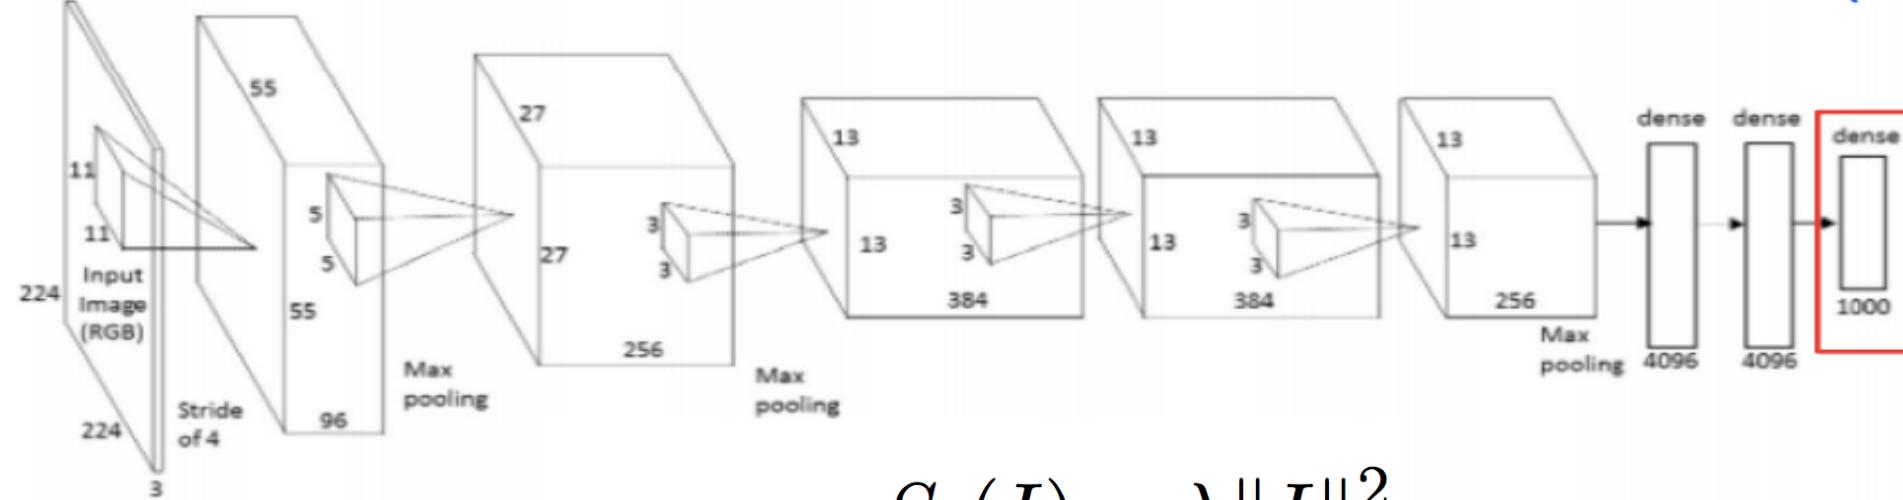


Outline - Visualization

- What is Visualization?
- Visualize patches that maximally activate neurons
- Visualize the weights
- Gradient based approaches
- Optimization based approach

Optimization Approach

Can we find an image that increases some class score?



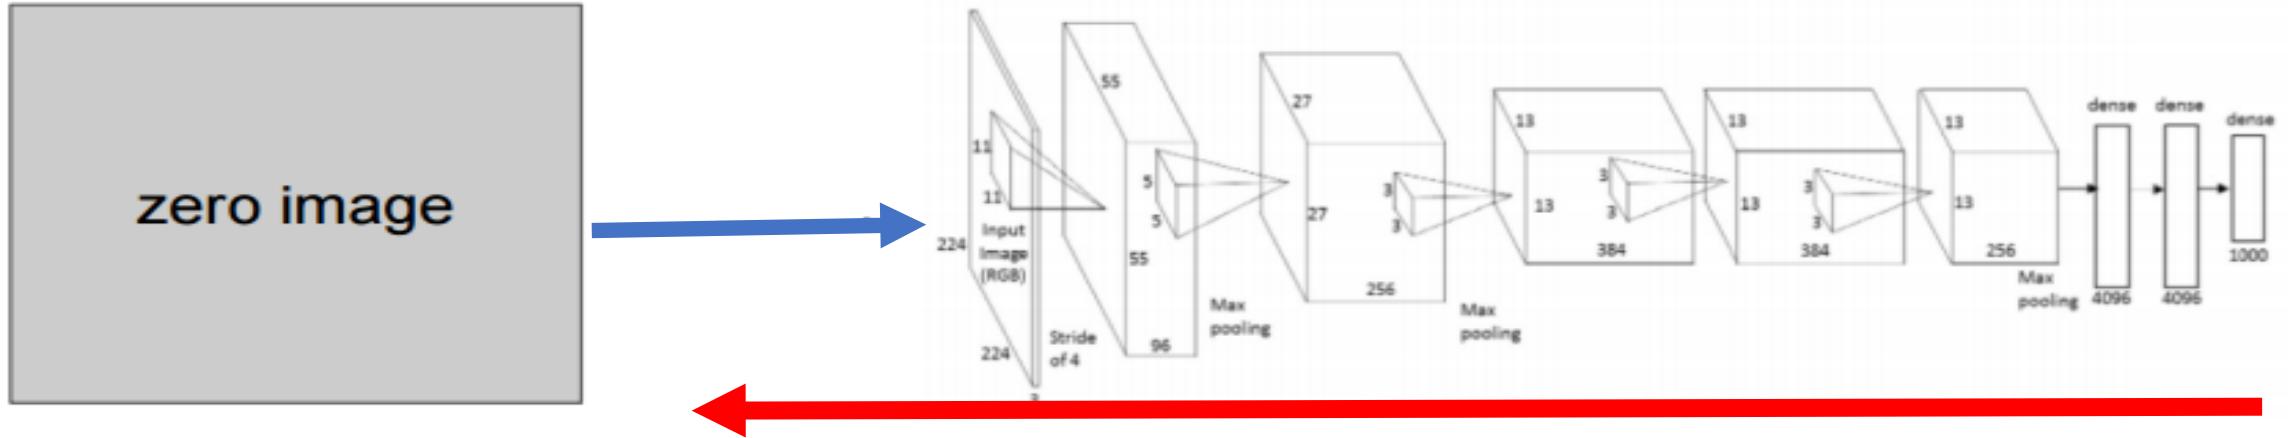
$$\arg \max_I S_c(I) - \lambda \|I\|_2^2$$

Score for class c before softmax

Regularization term



Optimization Approach - Algorithm



Start with zero image

Repeat:

Feed image forward

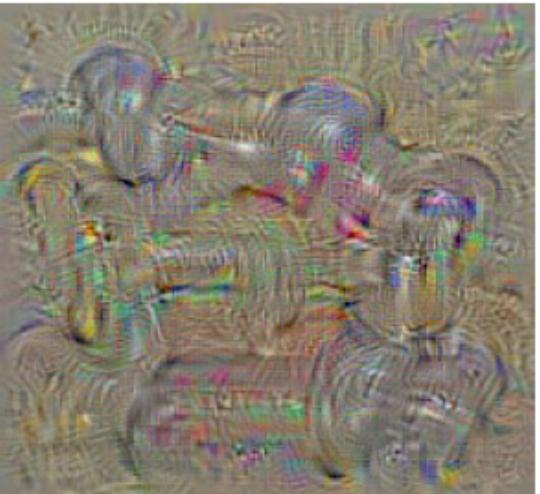
Set the gradient of the scores' vector to be [0,0,...,1,...,0]

Backward pass the gradients to the image

Update image (add regularization to avoid large updates)

$$\arg \max_I S_c(I) - \lambda \|I\|_2^2$$

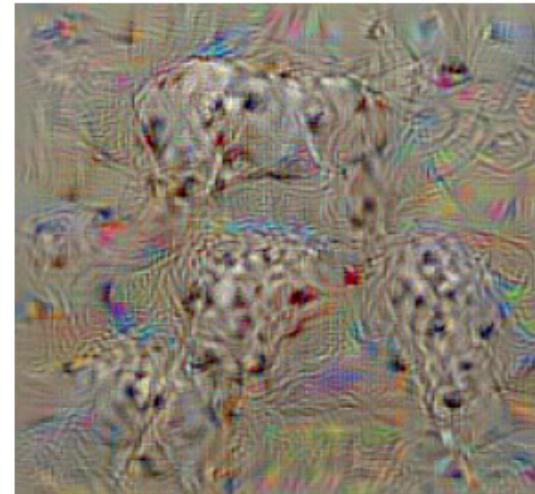
Optimization Approach - Examples



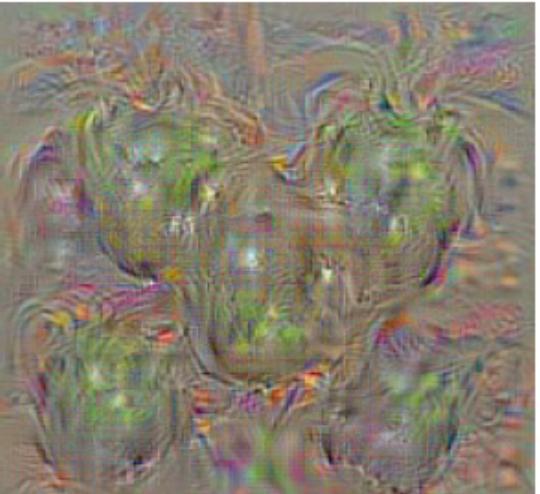
dumbbell



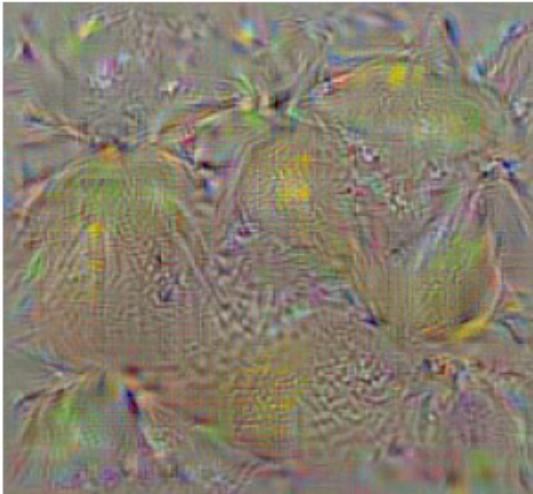
cup



dalmatian



bell pepper

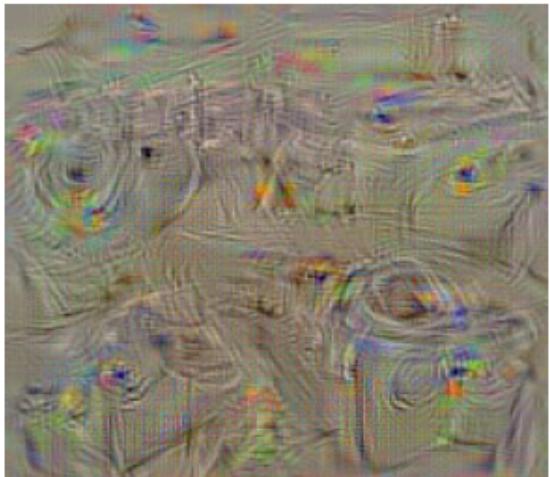


lemon

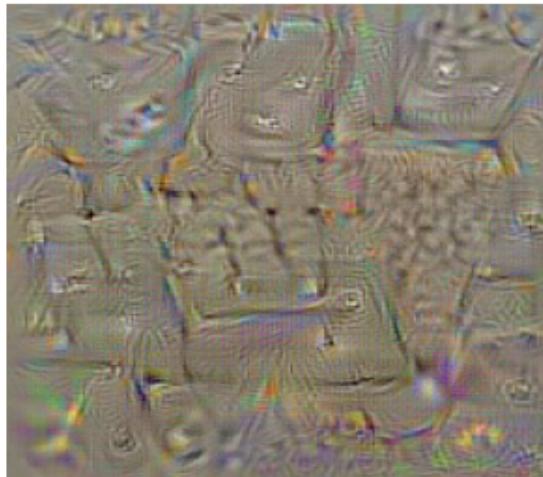


husky

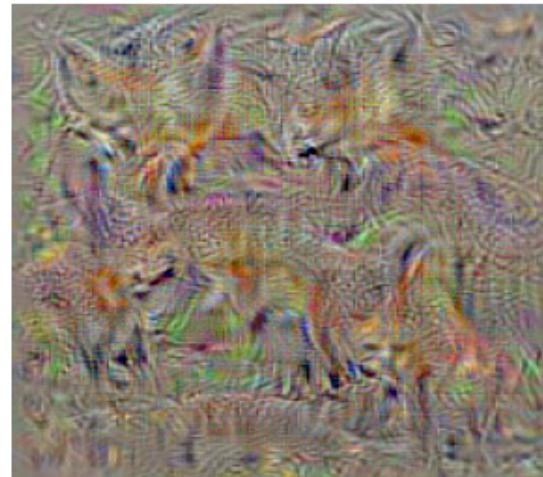
Optimization Approach - Examples



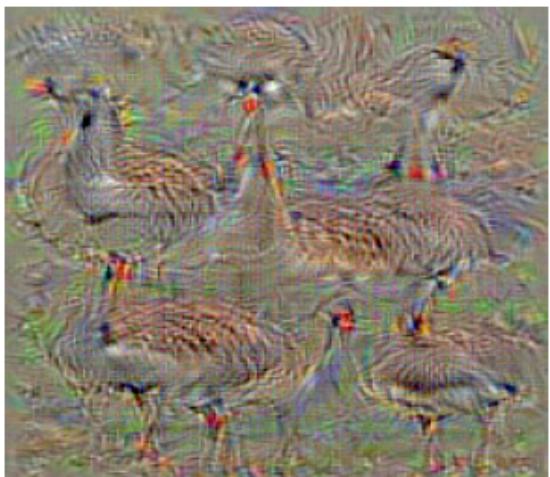
washing machine



computer keyboard



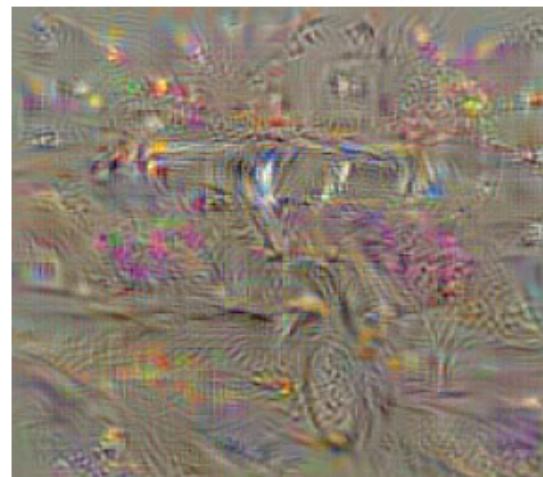
kit fox



goose

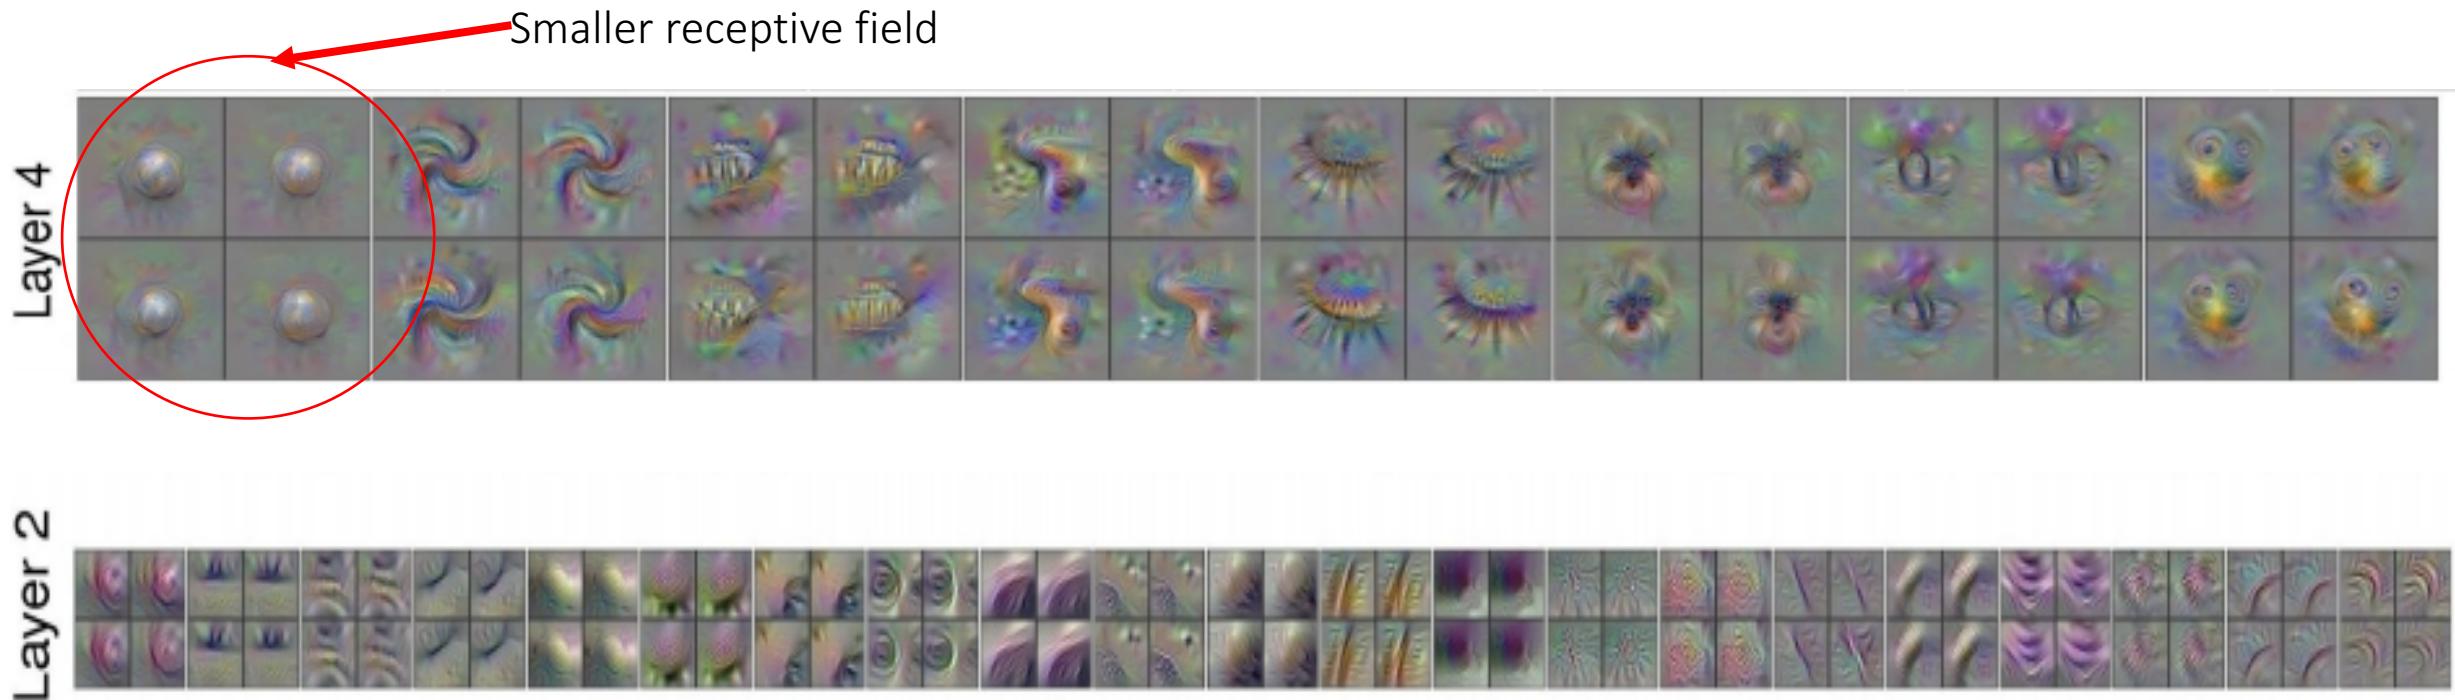


ostrich

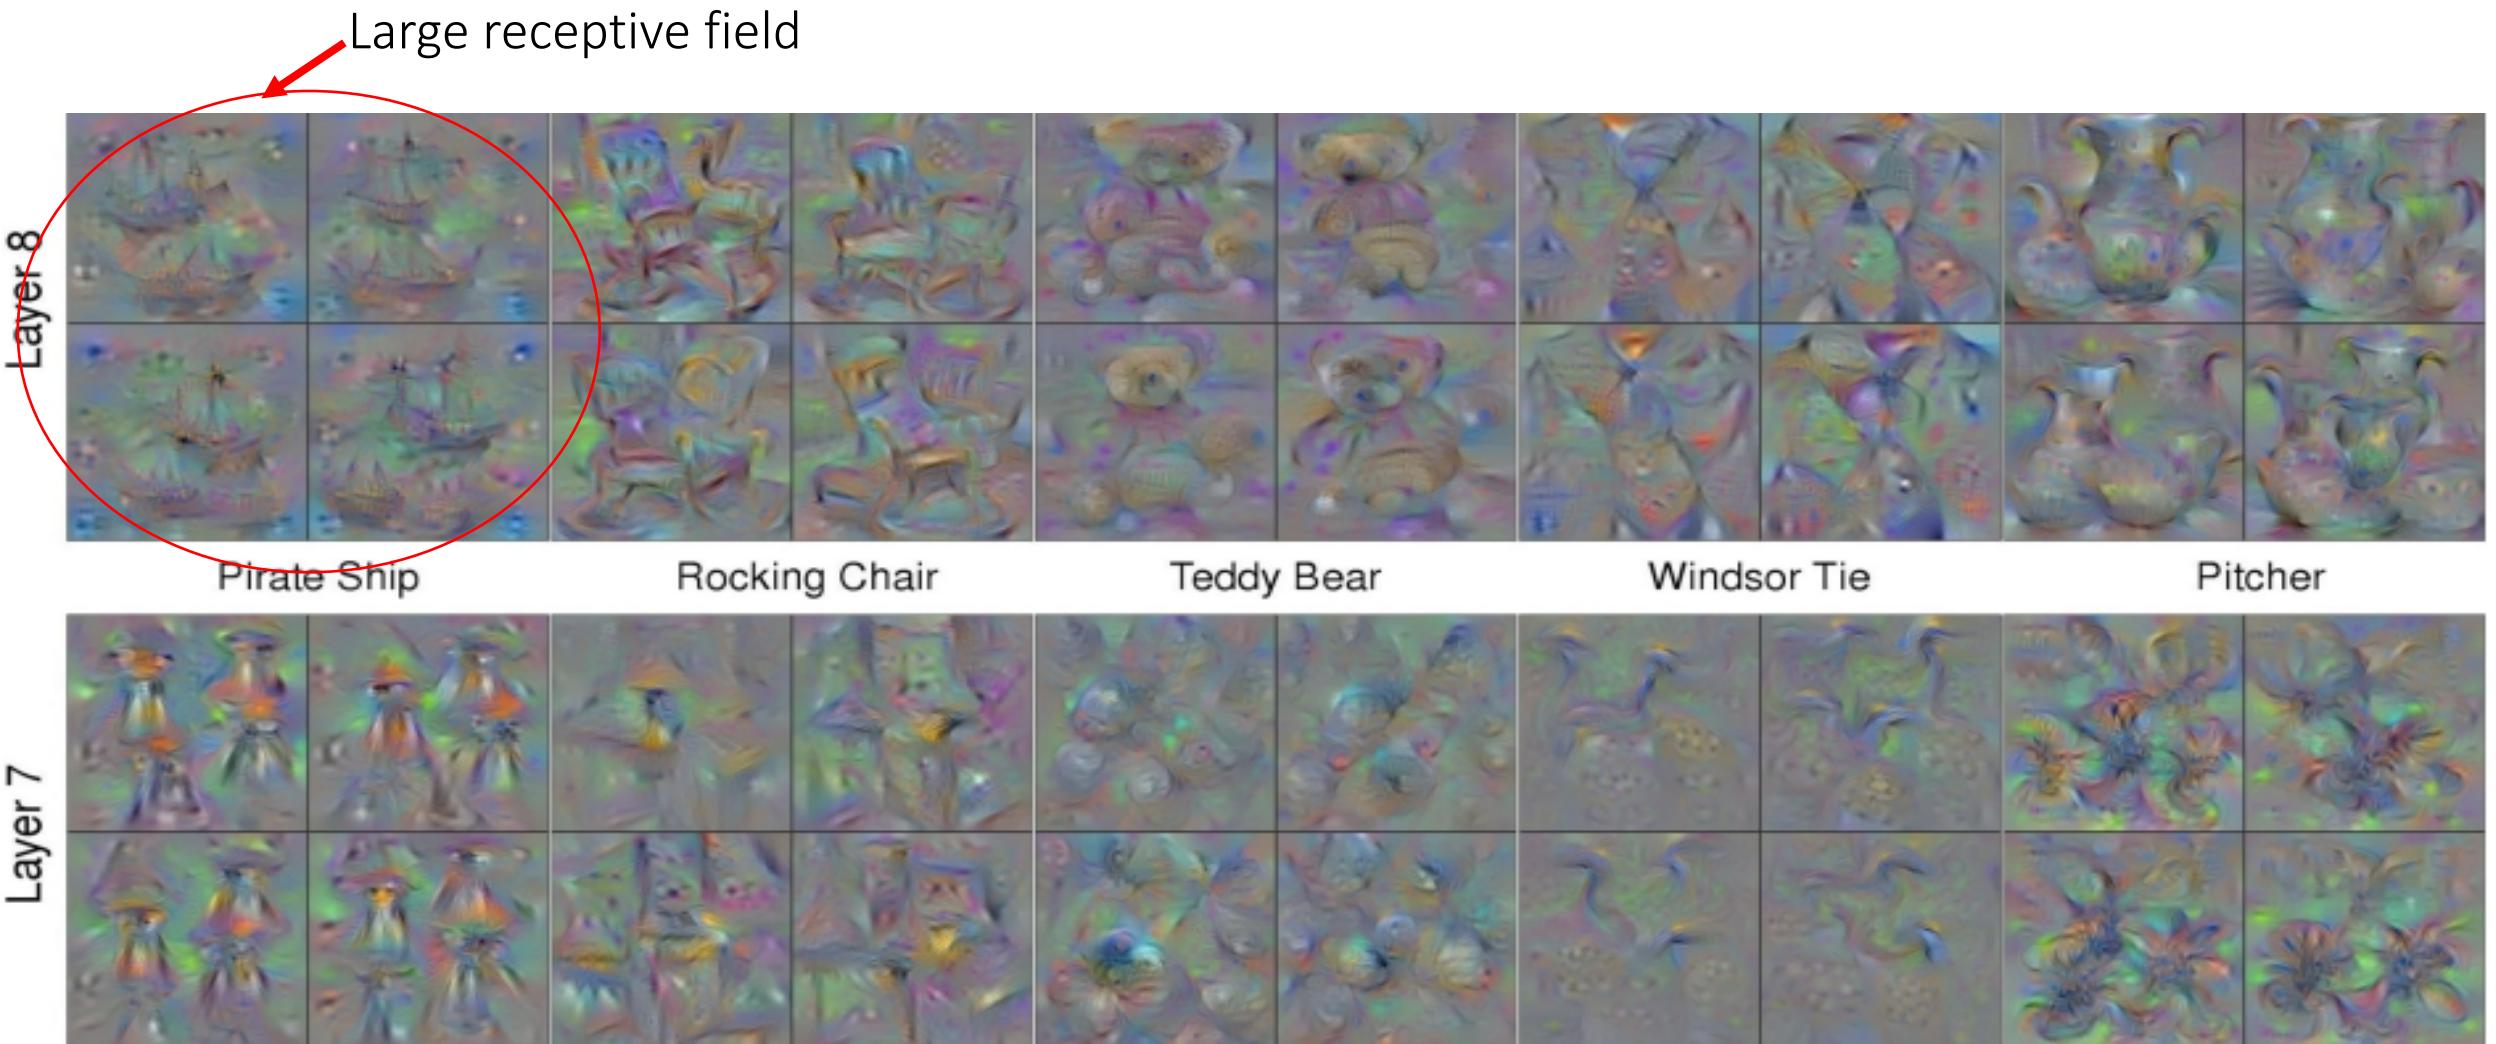


limousine

Visualizing Intermediate Layers

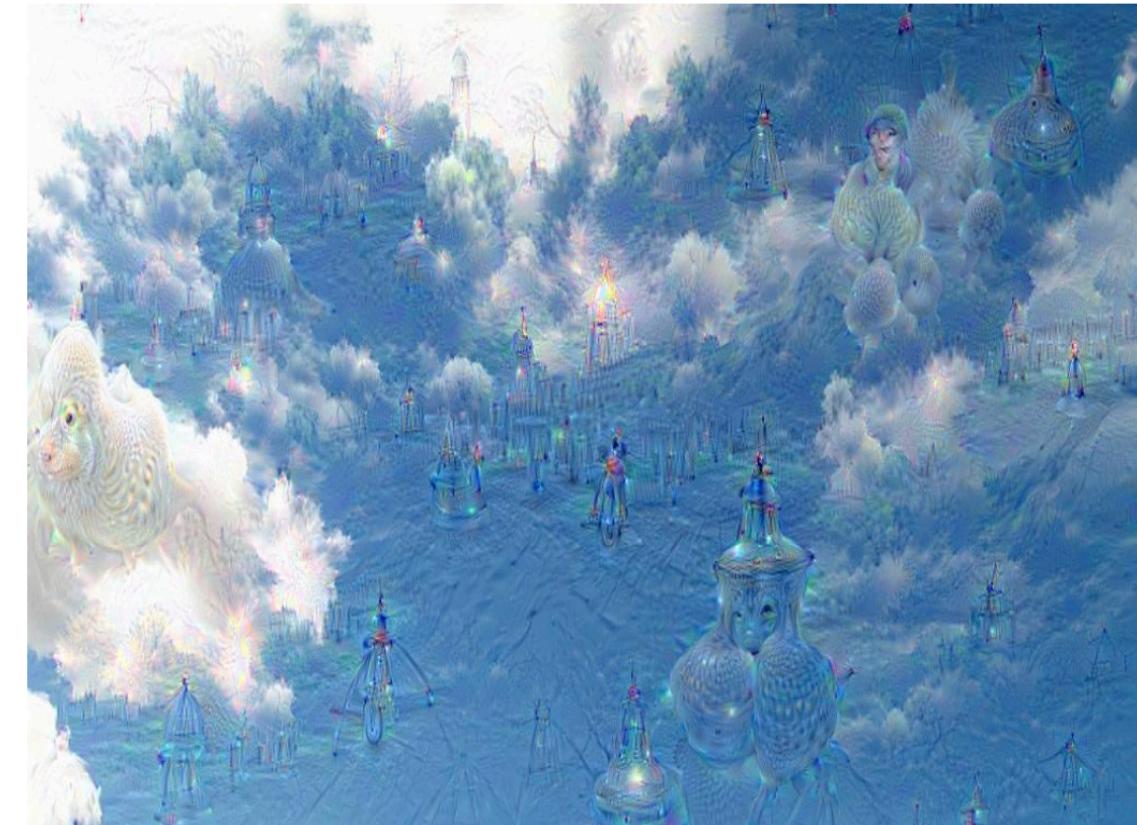


Visualizing Intermediate Layers

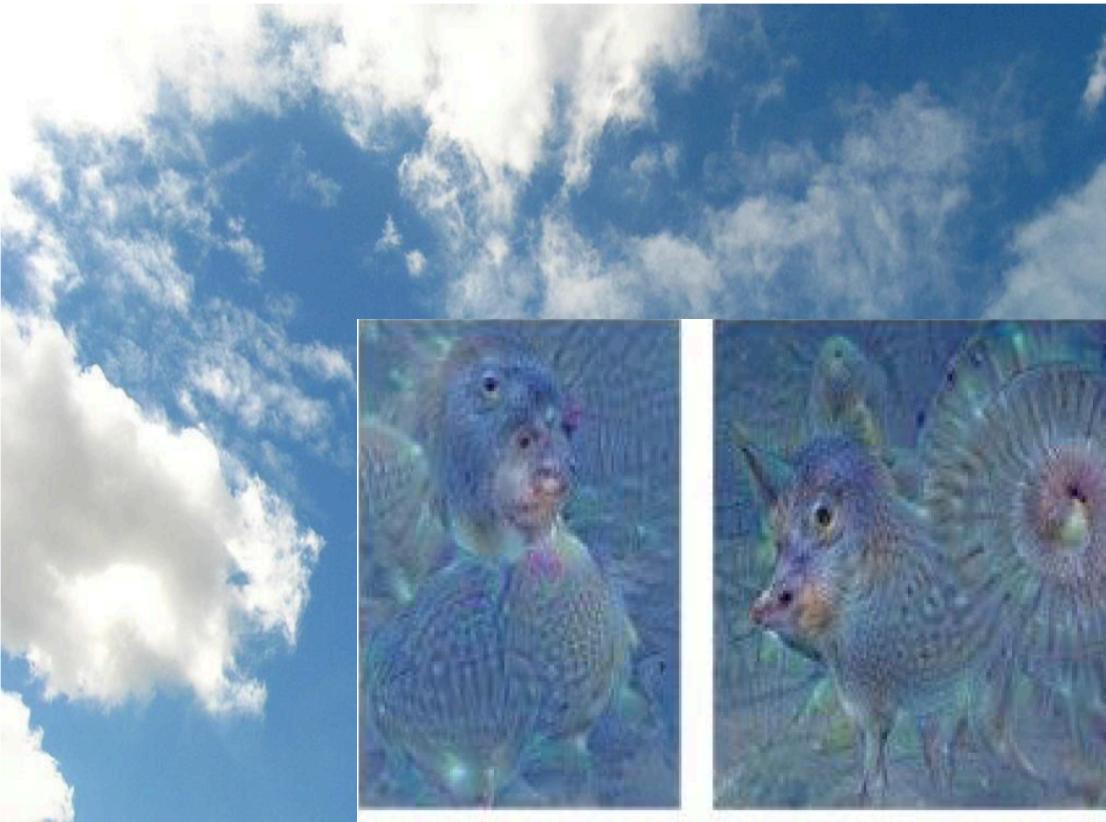


What if we map back the gradients onto the original image?

What if we map back the gradients onto the original image?



What if we map back the gradients onto the original image?



"Admiral Dog!"



"The Pig-Snail"



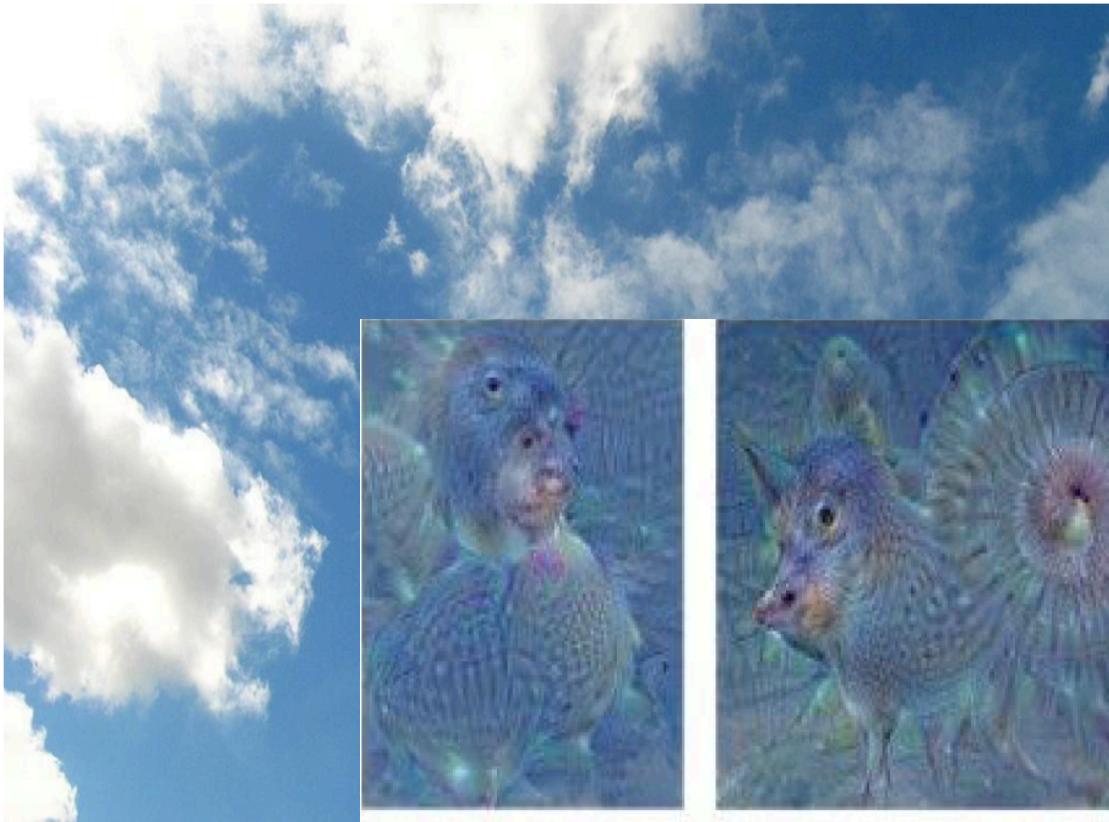
"The Camel-Bird"



"The Dog-Fish"

What if we map back the gradients onto the original image?

Deep Dream Grocery Store



"Admiral Dog!"



"The Pig-Snail"



"The Camel-Bird"

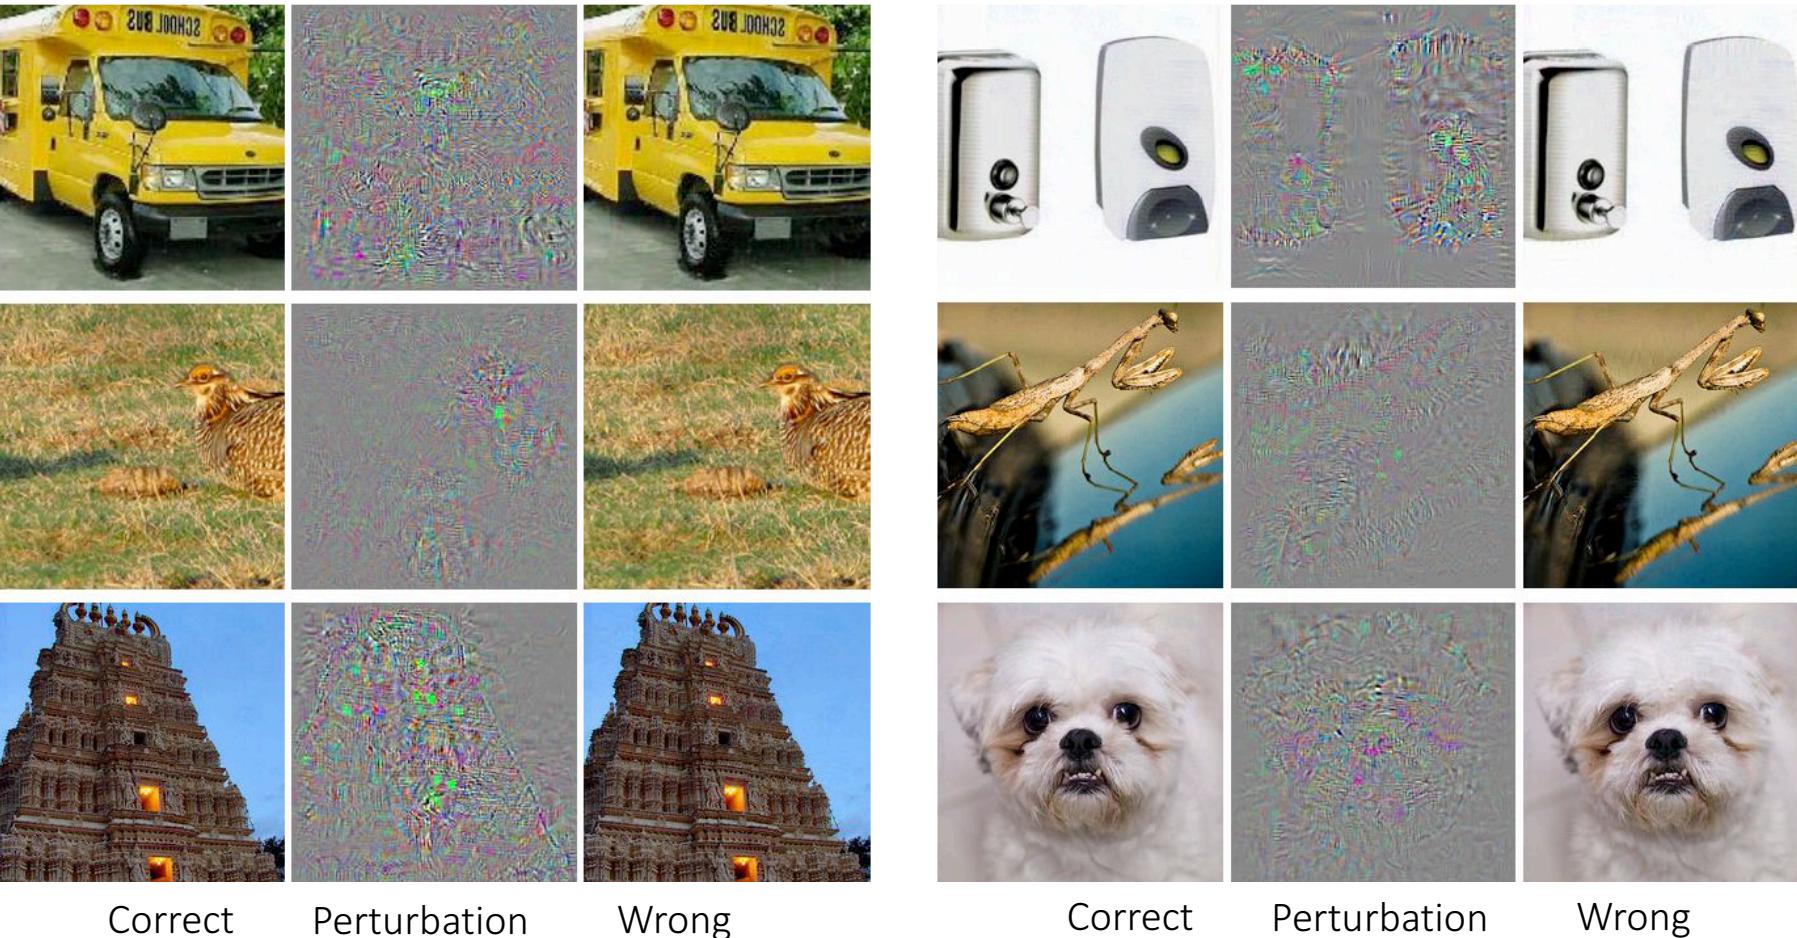


"The Dog-Fish"

Q: What is the difference between the gradient approach and the optimization approach for visualization?



Adversarial Examples

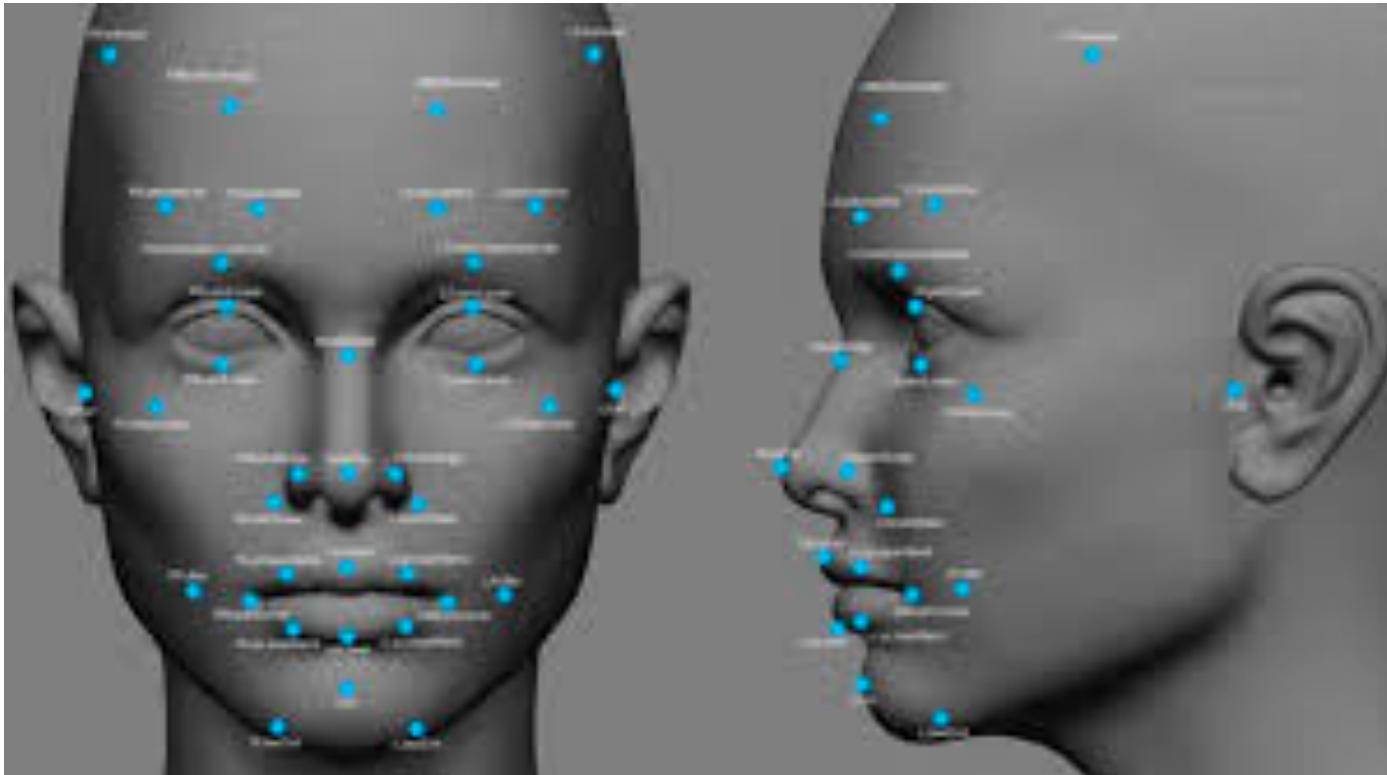


$$K(X + v) \neq K(X),$$

where K is a classifier, X is input image, v is perturbation.

Intriguing properties of neural networks, Szegedy et al. - 2013

Why care about adversarial examples?



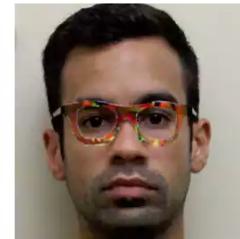
Why care about adversarial examples?



(b)



(c)

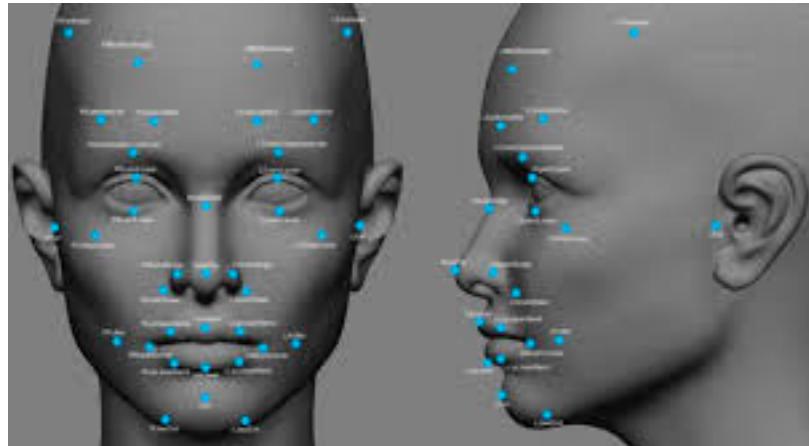


(d)

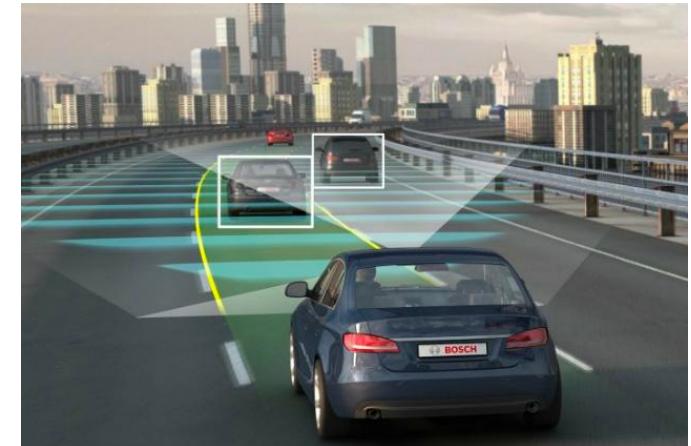


cvdazzle.com

Why care about adversarial examples?



Biometrics



Autonomous Driving



Security Guard Robot

“Build **safe**, widely distributed AI.”
-- OpenAI



Speech Recognition

Outline –Adversarial Examples

1. Adversarial and Rubbish examples
2. Evolutionary approach
3. Gradient based approaches
4. Adversarial training
5. Transferability
6. Universal Adversarial Perturbations
7. Why are neural networks easily fooled?
8. Proposed Solutions for adversarial attack

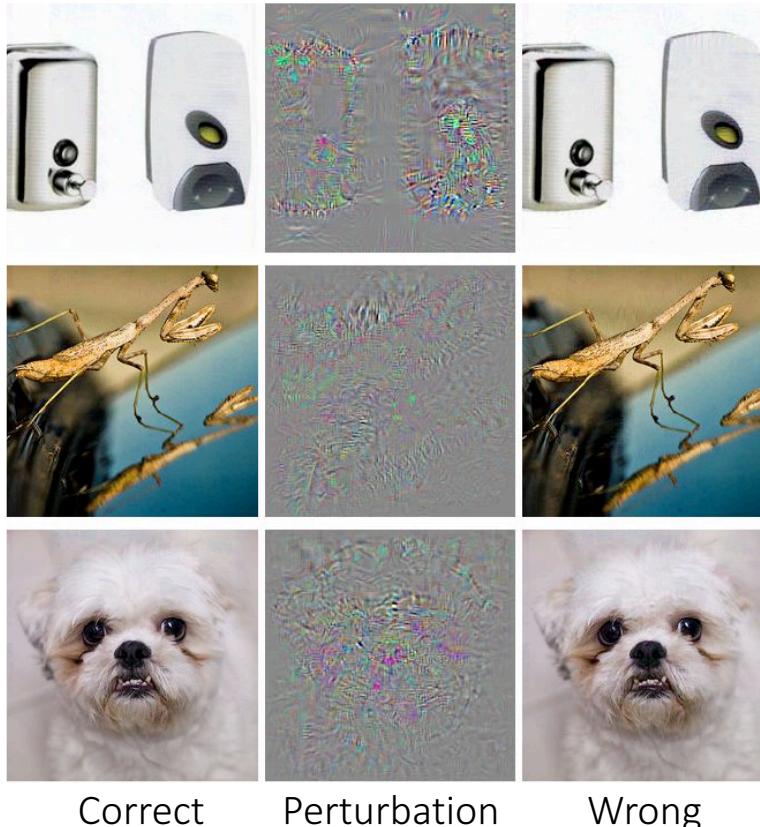
Outline –Adversarial Examples

1. Adversarial and Rubbish examples
2. Evolutionary approach
3. Gradient based approaches
4. Adversarial training
5. Transferability
6. Universal Adversarial Perturbations
7. Why are neural networks easily fooled?
8. Proposed Solutions for adversarial attack

Adversarial and Rubbish examples

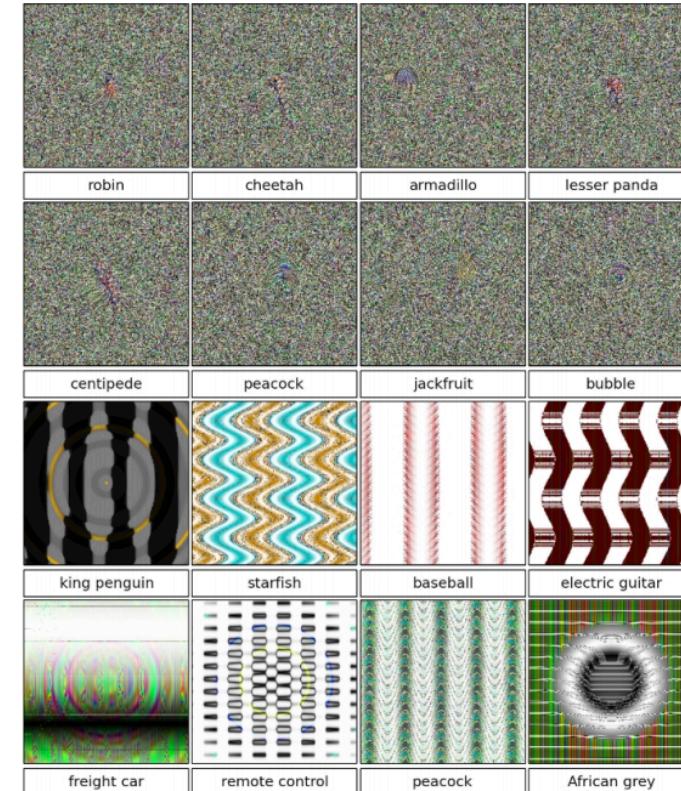
Adversarial

- corrupt an existing natural image



Rubbish

- noisy meaningless pictures that achieve high confidence classification



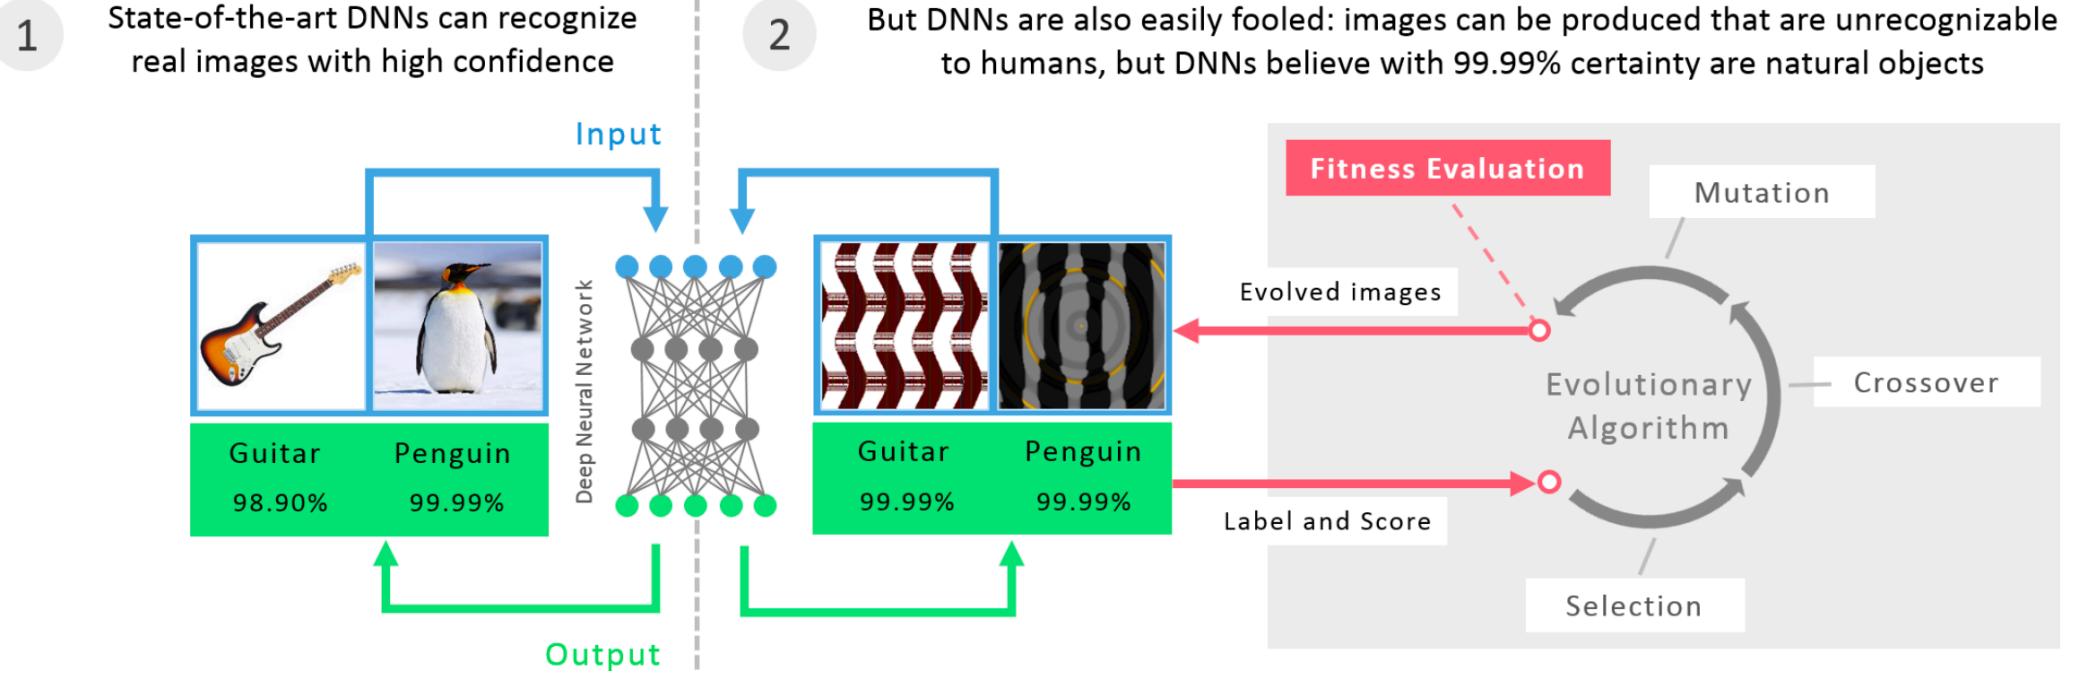
Intriguing properties of neural networks, Szegedy et al. - 2013

Deep Neural Networks are Easily Fooled: High Confidence Predictions for Unrecognizable Images - Nguyen, et al - 2014

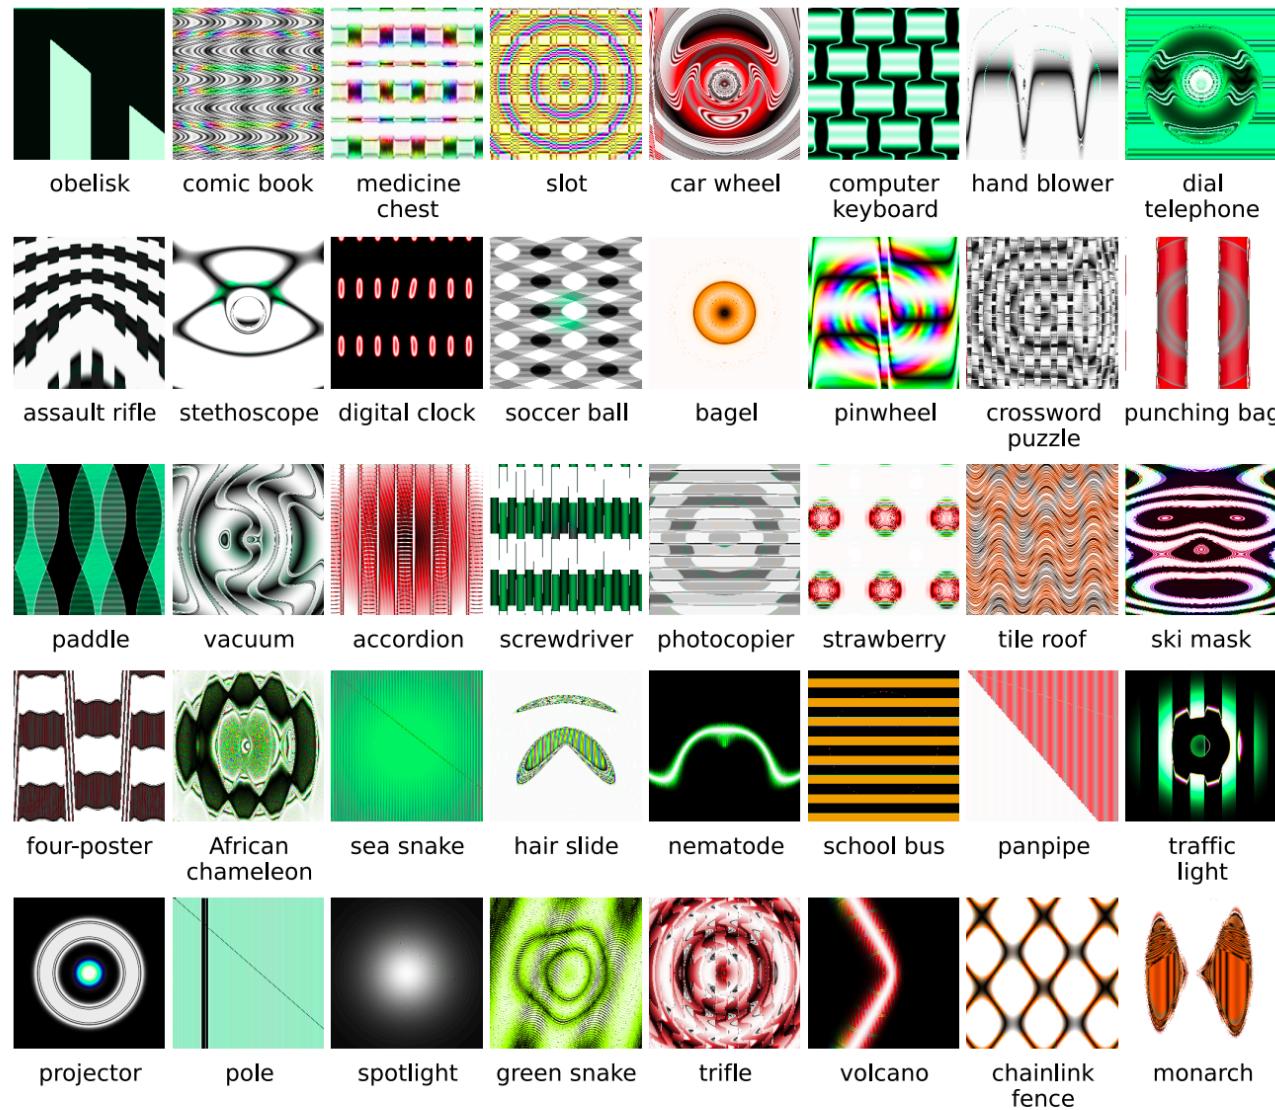
Outline –Adversarial Examples

1. Adversarial and Rubbish examples
2. Evolutionary approach
3. Gradient based approaches
4. Adversarial training
5. Transferability
6. Universal Adversarial Perturbations
7. Why are neural networks easily fooled?
8. Proposed Solutions for adversarial attack

Evolutionary Approach



Rubbish examples by evolutionary approach



Q: How can we change the image to fool the classifier?



Outline –Adversarial Examples

1. Adversarial and Rubbish examples
2. Evolutionary approach
- 3. Gradient based approaches**
4. Adversarial training
5. Transferability
6. Universal Adversarial Perturbations
7. Why are neural networks easily fooled?
8. Proposed Solutions for adversarial attack

Gradient-based approaches for visualization

$$\arg \max_I S_c(I) - \lambda \|I\|_2^2$$

Input image

The diagram illustrates the components of the optimization equation. Two red arrows point upwards from the text 'Score of class c given input image' and 'Regularization term' to the terms $S_c(I)$ and $\lambda \|I\|_2^2$ respectively. A blue arrow points downwards from the text 'Input image' to the variable I in the equation.

Score of class c given input image Regularization term

Gradient-based approaches for ~~visualization~~ adversarial examples

Visualization:

$$\arg \max_I S_c(I) - \lambda \|I\|_2^2$$

Adversarial examples:

1. Let $S_c(I)$ have high score for input I
2. We maximize the $-S_c(I + \text{noise})$ w.r.t noise
3. and penalize the L2-norm of noise.
4. We get a **new image $X = (I + \text{noise})$**

Fast Gradient Sign Method

Score of label y_{true} , given input image \mathbf{X}

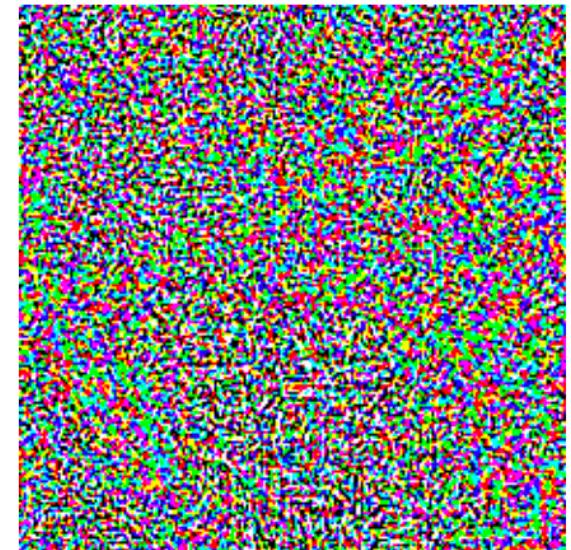
$$\mathbf{X}^{adv} = \mathbf{X} + \epsilon \text{sign}(\nabla_{\mathbf{X}} J(\mathbf{X}, y_{true}))$$



=



+ .007 ×



adversarial perturbation

Fast Gradient Sign Method

$$\mathbf{X}^{adv} = \mathbf{X} + \epsilon \text{sign}(\nabla_{\mathbf{X}} J(\mathbf{X}, y_{true}))$$



“gibbon”

=



“panda”

+ .007 ×



adversarial perturbation

Fast Gradient Sign Method

$$\mathbf{X}^{adv} = \mathbf{X} + \epsilon \text{sign}(\nabla_{\mathbf{X}} J(\mathbf{X}, y_{true}))$$



“gibbon”

=



“panda”

+ .007 ×



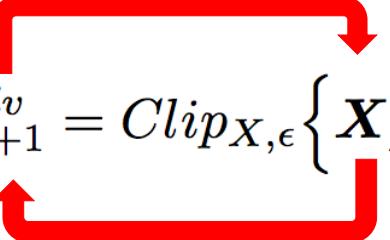
adversarial perturbation

Gradients-based Methods

- Fast Gradient Sign Method:

$$\mathbf{X}^{adv} = \mathbf{X} + \epsilon \text{sign}(\nabla_{\mathbf{X}} J(\mathbf{X}, y_{true}))$$

- Iterative Gradient Sign Method

$$\mathbf{X}_0^{adv} = \mathbf{X}, \quad \mathbf{X}_{N+1}^{adv} = Clip_{X,\epsilon} \left\{ \mathbf{X}_N^{adv} + \alpha \text{sign}(\nabla_{\mathbf{X}} J(\mathbf{X}_N^{adv}, y_{true})) \right\}$$


Iteratively repeat

Gradients-based Methods

- Fast Gradient Sign Method:

$$\mathbf{X}^{adv} = \mathbf{X} + \epsilon \operatorname{sign}(\nabla_{\mathbf{X}} J(\mathbf{X}, y_{true}))$$

- Iterative Gradient Sign Method

$$\mathbf{X}_0^{adv} = \mathbf{X}, \quad \mathbf{X}_{N+1}^{adv} = Clip_{X,\epsilon} \left\{ \mathbf{X}_N^{adv} + \alpha \operatorname{sign}(\nabla_{\mathbf{X}} J(\mathbf{X}_N^{adv}, y_{true})) \right\}$$

- Iterative Least-likely Class Method

$$y_{LL} = \arg \min_y \{ p(y|\mathbf{X}) \}$$

$$\mathbf{X}_0^{adv} = \mathbf{X}, \quad \mathbf{X}_{N+1}^{adv} = Clip_{X,\epsilon} \left\{ \mathbf{X}_N^{adv} - \alpha \operatorname{sign} (\nabla_{\mathbf{X}} J(\mathbf{X}_N^{adv}, y_{LL})) \right\}$$

Visual Comparison of Gradients-based Methods

Natural Image



Clean image



“Fast”; L_∞ distance to clean image = 32

Iterative Gradient Sign



“Basic iter.”; L_∞ distance to clean image = 32

Fast Gradient Sign



“L.l. class”; L_∞ distance to clean image = 28

Iterative LL-Class Gradient Sign

Outline –Adversarial Examples

1. Adversarial and Rubbish examples
2. Evolutionary approach
3. Gradient based approaches
- 4. Adversarial training**
5. Transferability
6. Universal Adversarial Perturbations
7. Why are neural networks easily fooled?
8. Proposed Solutions for adversarial attack

Adversarial Training

Q: How can we use adversarial examples to train a robust network?

A: Train it both on natural images and constructed adversarial images.

$$\tilde{J}(\theta, \mathbf{x}, y) = \alpha J(\theta, \mathbf{x}, y) + (1 - \alpha) J(\theta, \mathbf{x} + \epsilon \text{sign}(\nabla_{\mathbf{x}} J(\theta, \mathbf{x}, y)))$$

Training Target

Adversarial regularizer

Adversarial Training

How can we use adversarial examples to train a robust network?

$$\tilde{J}(\theta, \mathbf{x}, y) = \alpha J(\theta, \mathbf{x}, y) + (1 - \alpha) J(\theta, \mathbf{x} + \epsilon \text{sign}(\nabla_{\mathbf{x}} J(\theta, \mathbf{x}, y)))$$

Training Target Adversarial regularizer

For natural images, error rate drops from **0.94%** to **0.84%** on mnist.

For adversarial images, error rate drops from **89.4%** to **17.9%** on mnist.

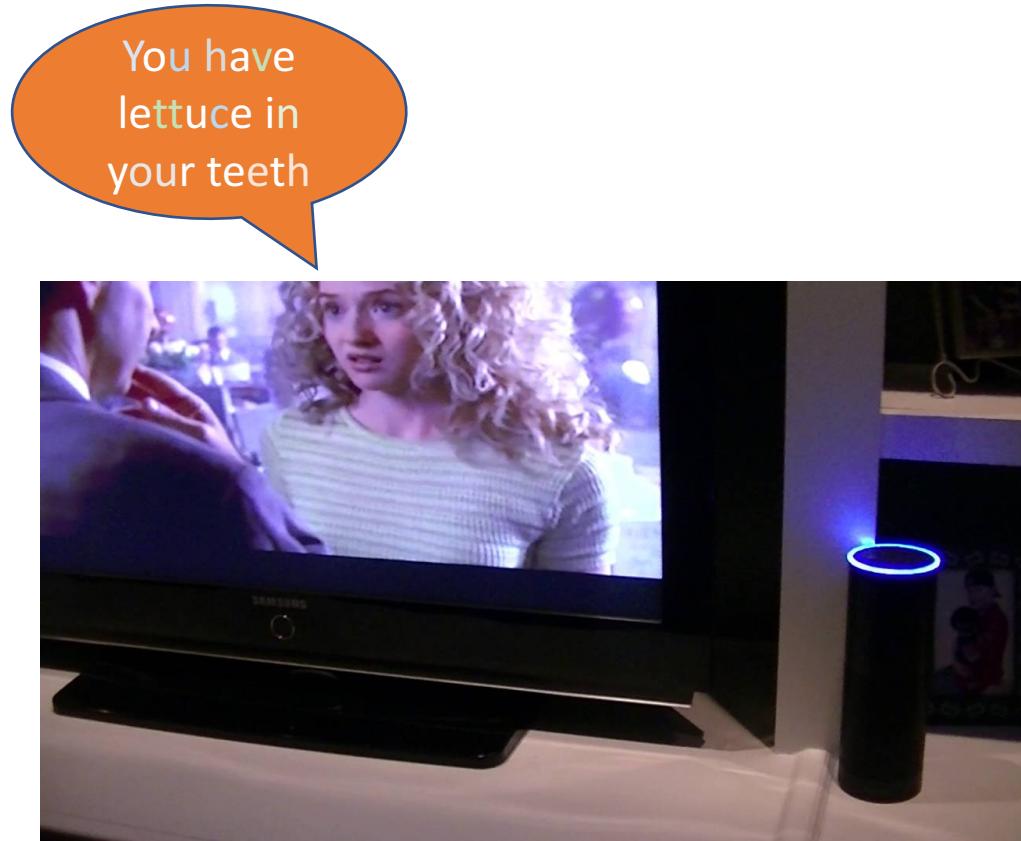
Outline – Adversarial Examples

1. Adversarial and Rubbish examples
2. Evolutionary approach
3. Gradient based approaches
4. Adversarial training
- 5. Transferability**
6. Universal Adversarial Perturbations
7. Why are neural networks easily fooled?
8. Proposed Solutions for adversarial attack

How much information do we need to fool a neural net?

Model weights	Have full access to model weights
Architecture	Know what the model looks like
Training data	Know what training dataset was used
Oracle/black box	Query model with input \mathbf{X} , get label \mathbf{Y}

Black box example – what we hear



Black box example – what we hear

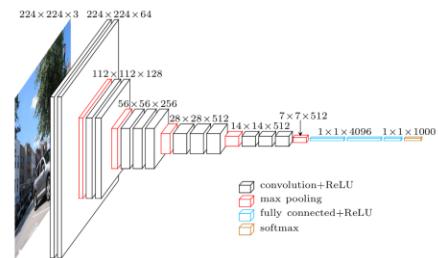


Transferability scenarios

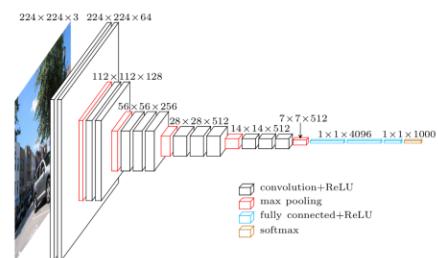
Cross training-set generalization

- Same architecture, different training set

8 2 9 4 4 6 4
1 3 5 9 1 7 6



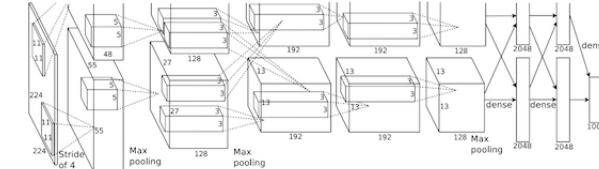
1 1 8 3 6 1 0
2 6 4 7 1 8 9



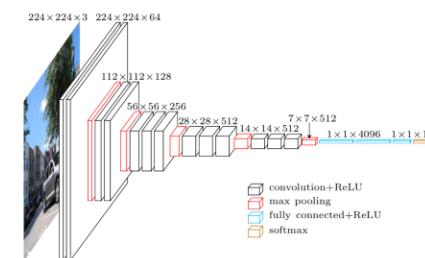
Cross model generalization

- Different architecture, same training set

IM₂GENET



IM₂GENET



Generalization error rates

	Model 1 - Dataset 1	Model 2 - Dataset 1	Model 1 - Dataset 2
M1 D1	100%	26.2%	5.9%
M2 D1	6.25%	100%	5.1%
M1 D2	8.2%	8.2%	100%
Gaussian noise	2.2%	2.6%	2.4%

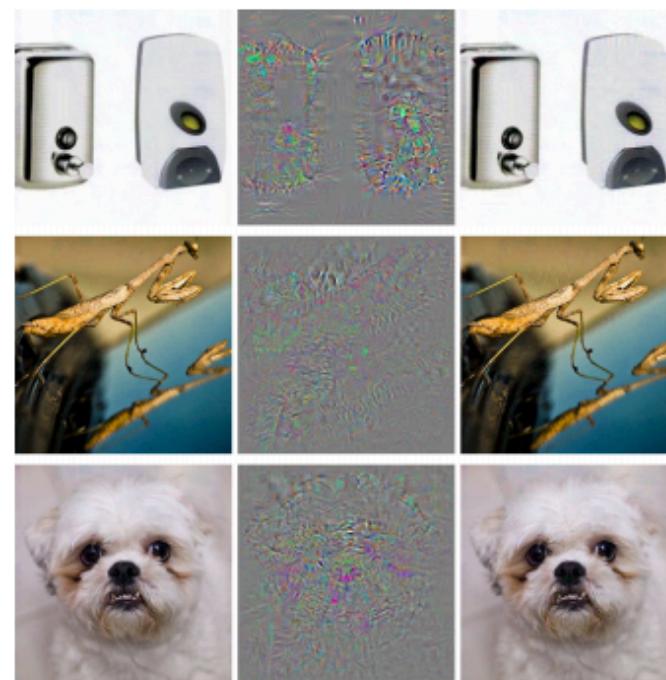
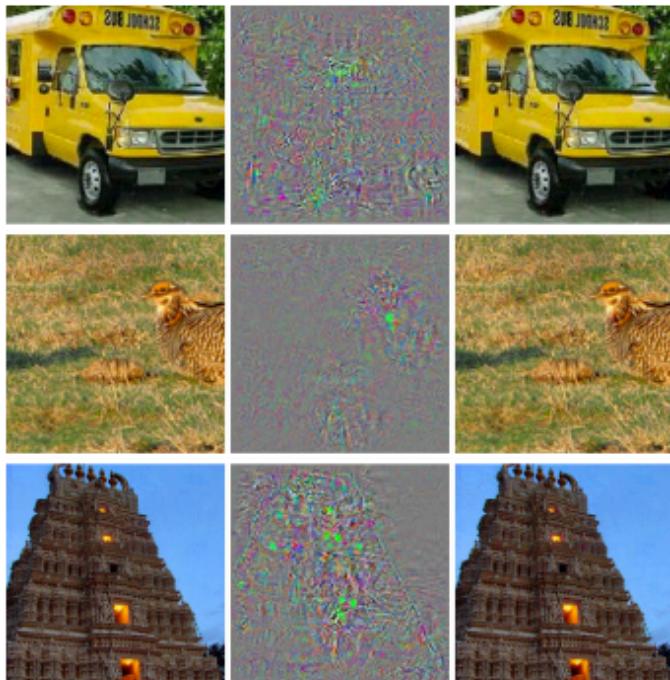
Table 1: Fooling rate for average perturbation $\text{stddev} = 0.06$

Generalization error rates

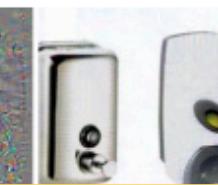
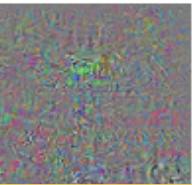
	Model 1 - Dataset 1	Model 2 - Dataset 1	Model 1 - Dataset 2
M1 D1	100%	98%	43%
M2 D1	96%	100%	22%
M1 D2	27%	50%	100%
Gaussian noise	2.6%	2.8%	2.7%

Table 2: Fooling rate for average perturbation `stddev = 0.1`

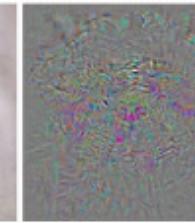
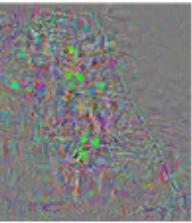
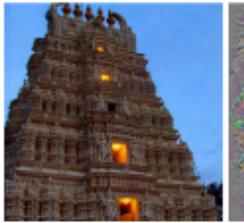
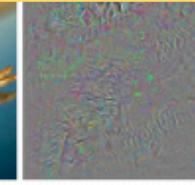
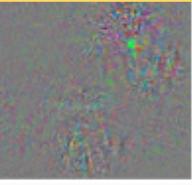
This is a very inefficient process



This is a very inefficient process



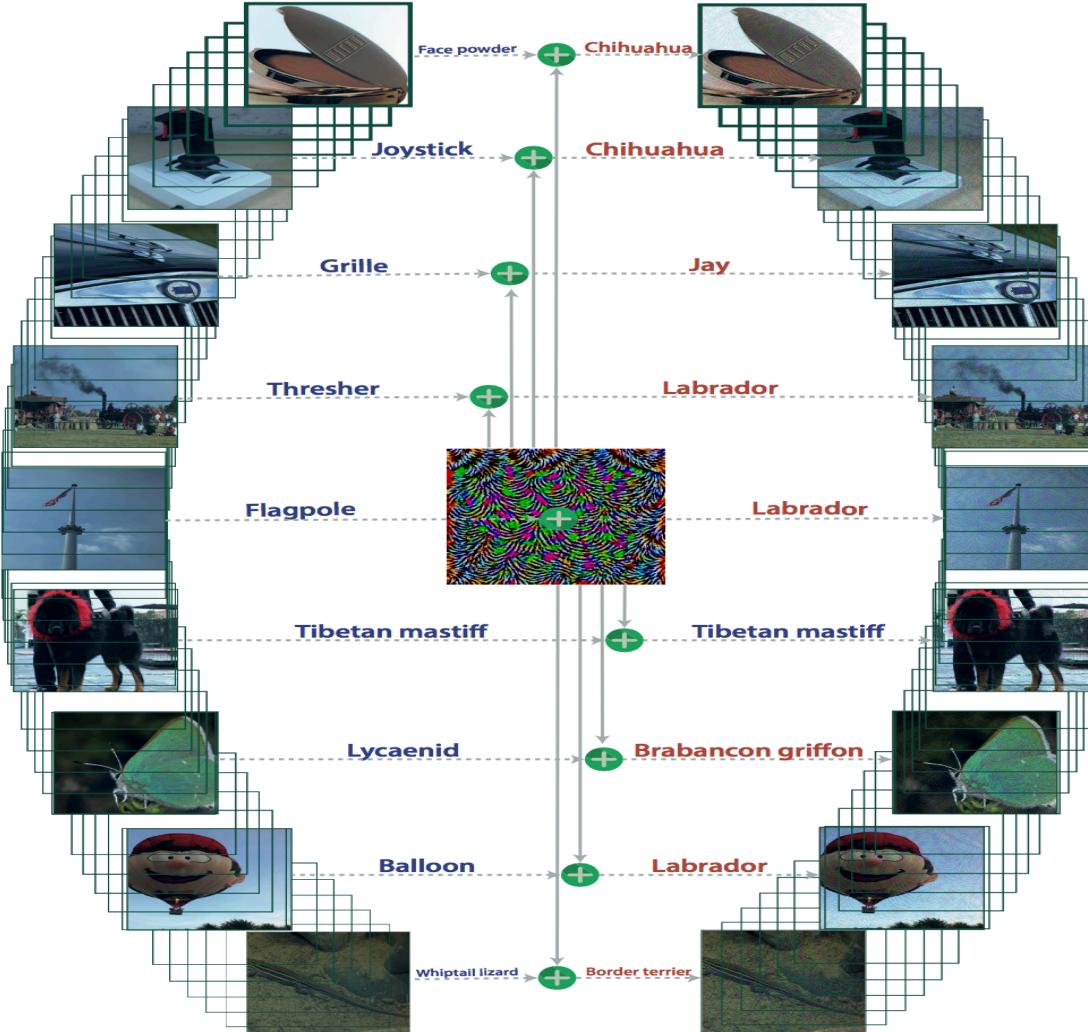
Q: what is the missing transferability property?



Outline –Adversarial Examples

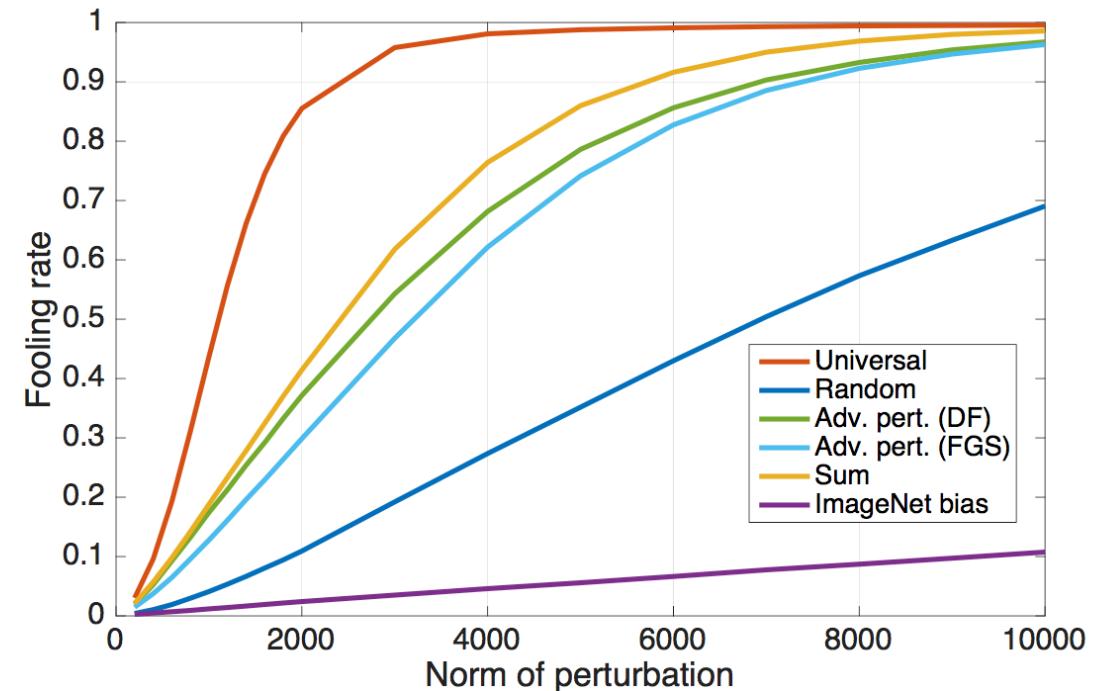
1. Adversarial and Rubbish examples
2. Evolutionary approach
3. Gradient based approaches
4. Adversarial training
5. Transferability
6. **Universal Adversarial Perturbations**
7. Why are neural networks easily fooled?
8. Proposed Solutions for adversarial attack

Universal Adversarial Perturbations



Candidate universal perturbations

- **Random noise**
 - Easy to compute
 - Needs high norm to be effective
 - Obvious to human
- **Sum of all adversarial perturbations over X**
 - Less obvious
 - Components known to be effective
 - Very expensive (compute $|X|$ times)
- **Universal Adversarial Perturbations (new method)**
 - Adaptively expensive (compute for a subset of X)
 - Very subtle



Algorithm

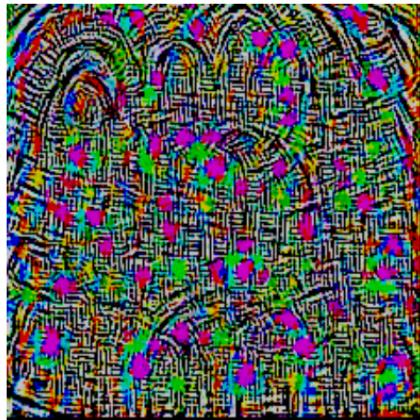
Intuition:

1. Start with $v = 0$
2. If $(X_i + v)$ is misclassified, skip to X_{i+1}
3. Find minimum perturbation Δv that takes $X_i + v + \Delta v$ to another class
4. Update $v = v + \Delta v$
5. Repeat with X_{i+1}

Algorithm 1 Computation of universal perturbations.

```
1: input: Data points  $X$ , classifier  $\hat{k}$ , desired  $\ell_p$  norm of  
   the perturbation  $\xi$ , desired accuracy on perturbed sam-  
   ples  $\delta$ .  
2: output: Universal perturbation vector  $v$ .  
3: Initialize  $v \leftarrow 0$ .  
4: while  $\text{Err}(X_v) \leq 1 - \delta$  do  
5:   for each datapoint  $x_i \in X$  do  
6:     if  $\hat{k}(x_i + v) = \hat{k}(x_i)$  then  
7:       Compute the minimal perturbation that  
       sends  $x_i + v$  to the decision boundary:  
         
$$\Delta v_i \leftarrow \arg \min_r \|r\|_2$$
  
         s.t.  $\hat{k}(x_i + v + r) \neq \hat{k}(x_i)$ .  
8:   Update the perturbation:  
   
$$v \leftarrow \mathcal{P}_{p,\xi}(v + \Delta v_i).$$
  
9:   end if  
10:  end for  
11: end while
```

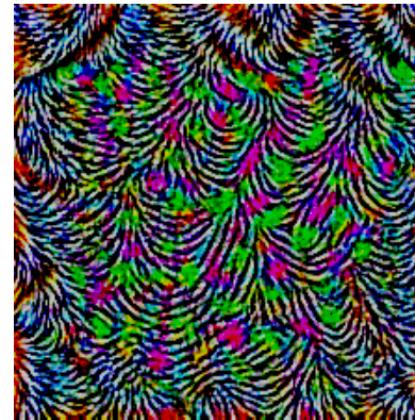
Sample universal perturbations



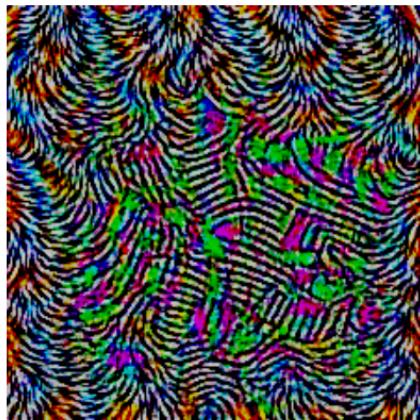
(a) CaffeNet



(b) VGG-F



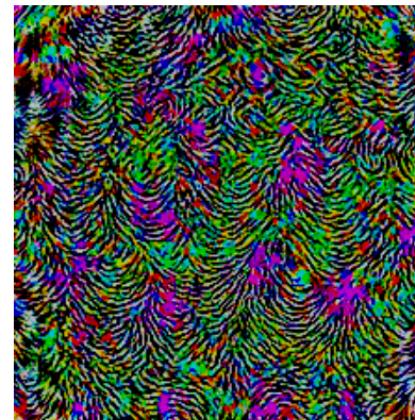
(c) VGG-16



(d) VGG-19



(e) GoogLeNet



(f) ResNet-152

Cross-model universality

	VGG-F	CaffeNet	GoogLeNet	VGG-16	VGG-19	ResNet-152
VGG-F	93.7%	71.8%	48.4%	42.1%	42.1%	47.4 %
CaffeNet	74.0%	93.3%	47.7%	39.9%	39.9%	48.0%
GoogLeNet	46.2%	43.8%	78.9%	39.2%	39.8%	45.5%
VGG-16	63.4%	55.8%	56.5%	78.3%	73.1%	63.4%
VGG-19	64.0%	57.2%	53.6%	73.5%	77.8%	58.0%
ResNet-152	46.3%	46.3%	50.5%	47.0%	45.5%	84.0%

Fooling rate when computing a perturbation for one model (**rows**) and testing it on others (**columns**)

Outline – Adversarial Examples

1. Adversarial and Rubbish examples
2. Evolutionary approach
3. Gradient based approaches
4. Adversarial training
5. Transferability
6. Universal Adversarial Perturbations
7. **Why are neural networks easily fooled?**
8. Proposed Solutions for adversarial attack

Models are too linear

X	2	-1	3	-2	2	2	1	-4	5	1	← input example
w	-1	-1	1	-1	1	-1	1	1	-1	1	← weights

class 1 score = dot product:

$$= -2 + 1 + 3 + 2 + 2 - 2 + 1 - 4 - 5 + 1 = -3$$

Models are too linear

X	2	-1	3	-2	2	2	1	-4	5	1
W	-1	-1	1	-1	1	-1	1	1	-1	1
adversarial x	?	?	?	?	?	?	?	?	?	?

← input example
← weights

class 1 score = dot product:

$$= -2 + 1 + 3 + 2 + 2 - 2 + 1 - 4 - 5 + 1 = -3$$

Models are too linear

X	2	-1	3	-2	2	2	1	-4	5	1
W	-1	-1	1	-1	1	-1	1	1	-1	1
adversarial x	1.5	-1.5	3.5	-2.5	2.5	1.5	1.5	-3.5	4.5	1.5

← input example
← weights

class 1 score before:

$$-2 + 1 + 3 + 2 + 2 - 2 + 1 - 4 - 5 + 1 = -3$$

$$\textcolor{red}{-1.5+1.5+3.5+2.5-1.5+1.5-3.5-4.5+1.5 = 2}$$

Outline – Adversarial Examples

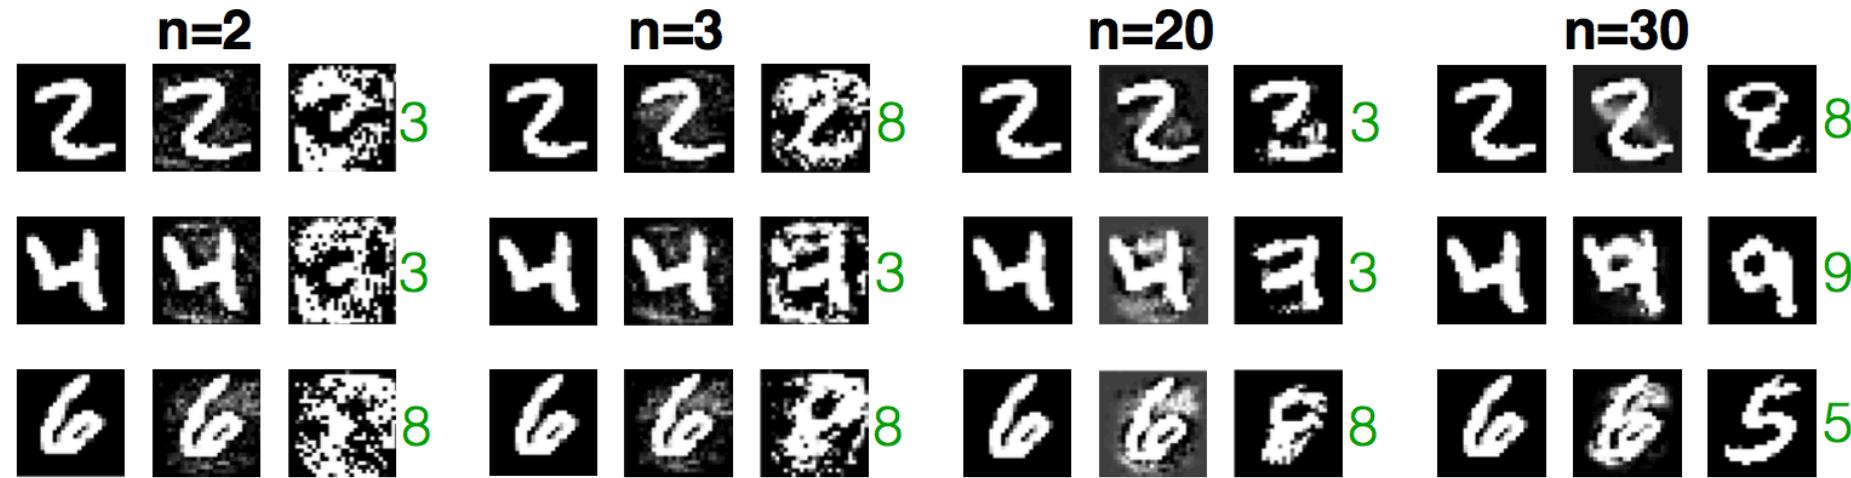
1. Adversarial and Rubbish examples
2. Evolutionary approach
3. Gradient based approaches
4. Adversarial training
5. Transferability
6. Universal Adversarial Perturbations
7. Why are neural networks easily fooled?
8. Proposed Solutions for adversarial attack

Proposed solution: highly non-linear models

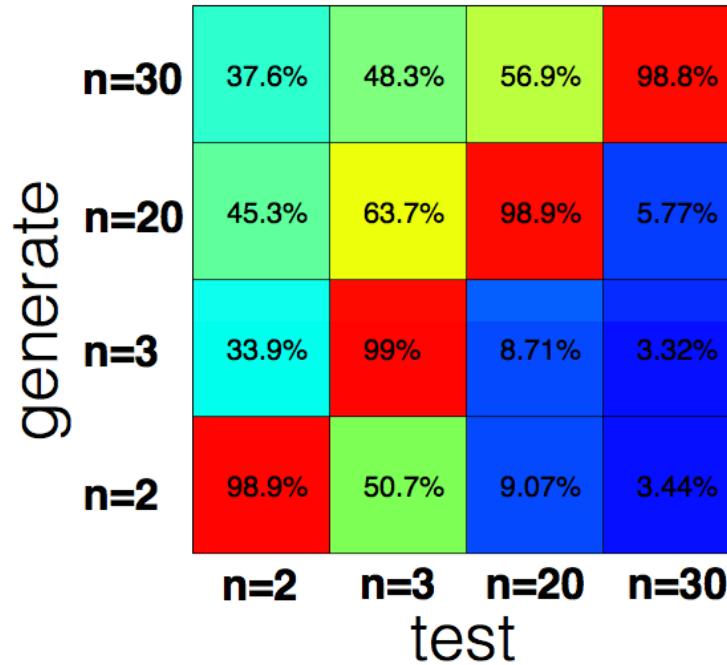
- Use a **rectified polynomial** as the activation

$$F_n(x) = \begin{cases} x^n, & x \geq 0 \\ 0, & x < 0 \end{cases}$$

Robustness against Adversarial Examples



Fooling Rate



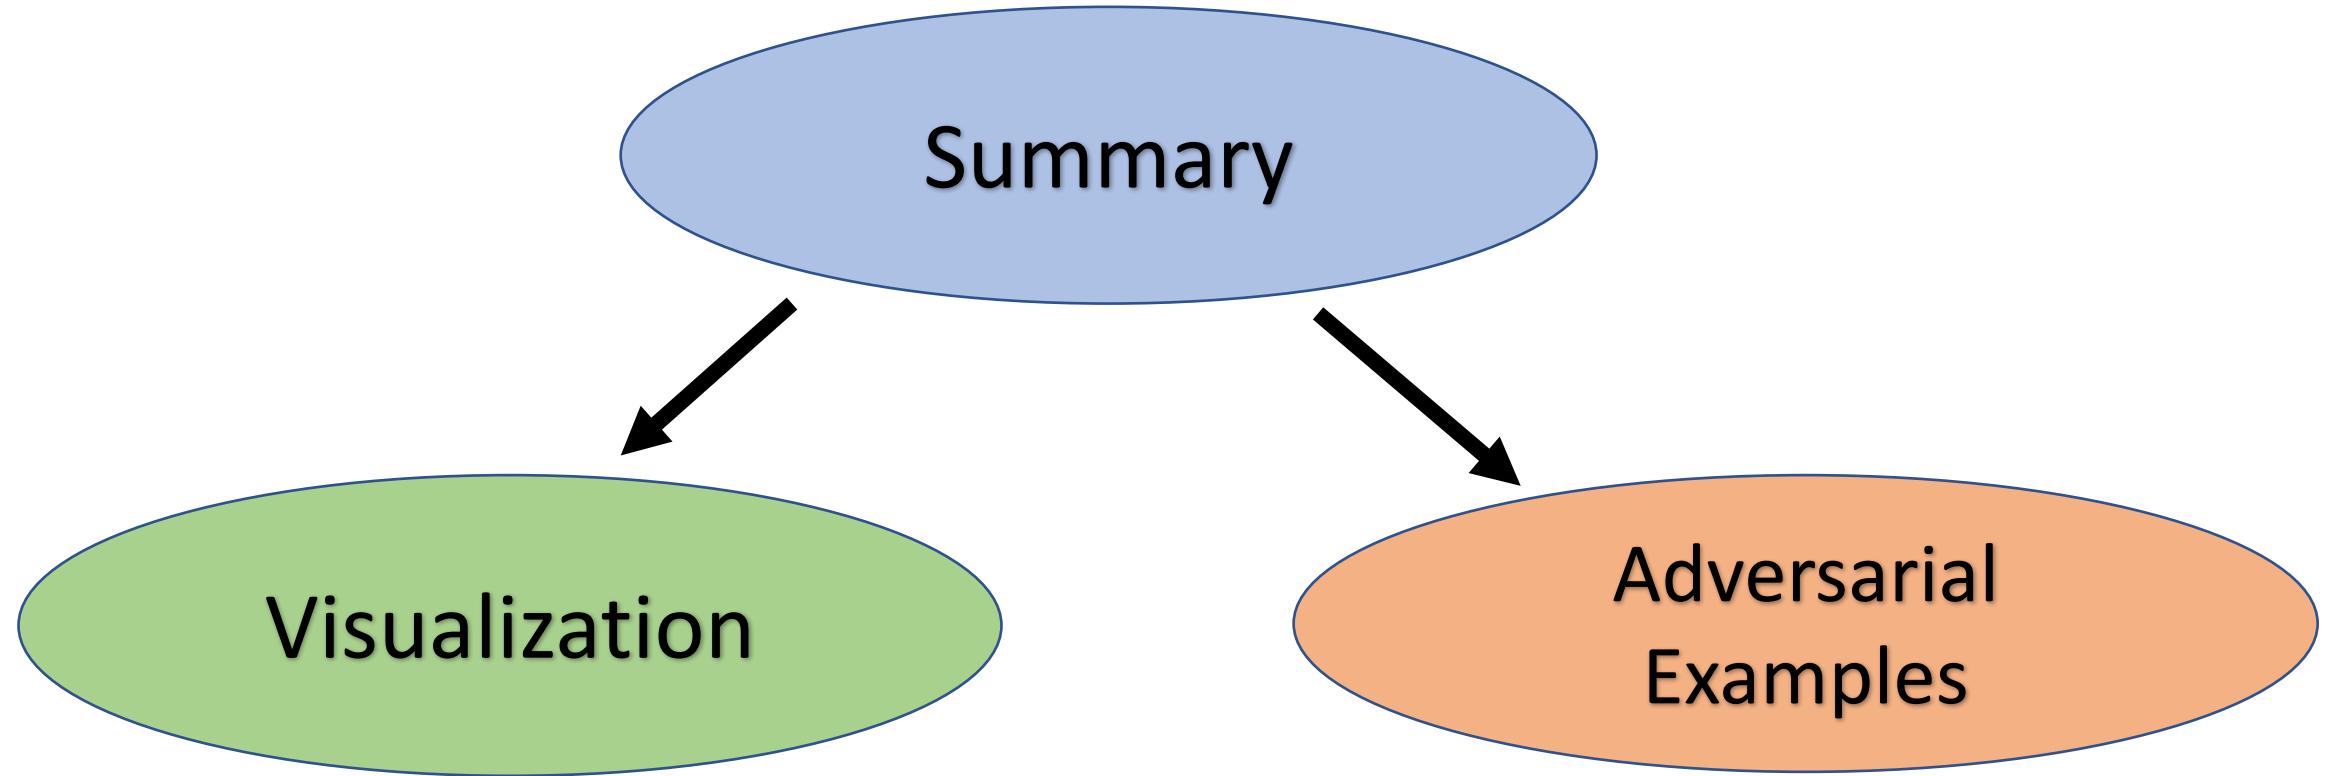
clean MNIST test set:

$$\text{error}_{n=2} = 1.51\%$$

$$\text{error}_{n=3} = 1.44\%$$

$$\text{error}_{n=20} = 1.61\%$$

$$\text{error}_{n=30} = 1.80\%$$



- ✓ What is Visualization?
- ✓ Visualize patches that maximally activate neurons
- ✓ Visualize the weights
- ✓ Gradient based approaches
- ✓ Optimization based approach

- ✓ Adversarial and Rubbish examples
- ✓ Evolutionary approach
- ✓ Gradient based approaches
- ✓ Adversarial training
- ✓ Transferability
- ✓ Universal Adversarial Perturbations
- ✓ Why are neural networks easily fooled?
- ✓ Proposed Solutions for adversarial attack

Reading list

- Matthew D. Zeiler, Rob Fergus [Visualizing and Understanding Convolutional Networks](#), ECCV 2014
- Karen Simonyan, Andrea Vedaldi, Andrew Zisserman, [Deep Inside Convolutional Networks: Visualising Image Classification Models and Saliency Maps](#) arXiv:1312.6034v2
- Alexey Dosovitskiy Thomas Brox, [Inverting Visual Representations with Convolutional Networks](#), CVPR 2016
- Anh Nguyen, Jason Yosinski, Jeff Clune. [Deep Neural Networks are Easily Fooled: High Confidence Predictions for Unrecognizable Images](#), CVPR 2015
- Christian Szegedy, et al. [Intriguing properties of neural networks](#), arXiv preprint arXiv:1312.6199v4
- Alexey Kurakin, et al, [Adversarial examples in the physical world](#), arXiv preprint arXiv:1607.02533
- Seyed-Mohsen Moosavi-Dezfooli, et al, [Universal adversarial perturbations](#), arXiv preprint arXiv:1610.08401v2
- Dmitry Krotov, et al, [Dense Associative Memory is Robust to Adversarial Inputs](#), arXiv preprint arXiv:1701.00939
- Ian J. Goodfellow, et al, [Explaining and Harnessing Adversarial Examples](#), arXiv preprint arXiv:1412.6572
- Nicholas Carlini et al, [Hidden Voice Commands](#), 25th USENIX Security Symposium
- Brian Chu et al, [Visualizing Residual Networks](#), arXiv preprint arXiv:1701.02362
- Nicolas Papernot et al, [SoK: Towards the Science of Security and Privacy in Machine Learning](#), arXiv preprint arXiv:1611.03814
- Nicolas Papernot et al, [Practical Black-Box Attacks against Deep Learning Systems using Adversarial Examples](#), arXiv preprint arXiv:1602.02697
- Ian J. Goodfellow et al, [Attacking machine learning with adversarial examples](#), OpenAI blog post

Conclusion

